

Assignment3_answers

Hair Parra

April 16, 2020

Hair ALbeiro Parra Barrera

260738619

MATH 545: ASSG3

Assignment 3

Libraries

```
library(tidyverse)
library(tidyquant)
library(gridExtra)
library(tibbletime)
library(forecast)
library(itsmr)
library(here)
library(bbmle)
library(tseries)
library(fpp2)
knitr::opts_chunk$set(comment=NA, tidy=FALSE)
#library(future) Not needed yet
#library(doFuture) Not needed yet
#library(rbenchmark) Not needed yet
```

Question 3.9

- (a) Calculate the autocovariance function $\gamma(\cdot)$ of the stationary time series:

$$Y_t = \mu_t + Z_t + \theta_1 Z_1 + \theta_{12} Z_{12}, \quad \{Z_t\} \sim WN(0, \sigma^2)$$

Solution

In the hand-written part , we got that

$$\gamma_Y(h) = \begin{cases} \sigma^2(1 + \theta_1^2 + \theta_{12}^2), & h = 0 \\ \sigma^2\theta_1, & h = 1 \\ \sigma^2\theta_1\theta_{12}, & h = 11 \\ \sigma^2\theta_{12}, & h = 12 \end{cases}$$

- (b) Compute the sample mean and sample autocovariances $\hat{\gamma}(h), 0 \leq h \leq 20$ of $\{\nabla\nabla_{12}X_t\}$, where $\{X_t, t = 1, \dots, 72\}$ is the accidental deaths of the sereis DEATHS

```
library(itsmr)
deaths <- ts(deaths) # convert into a time series
deaths
```

Time Series:

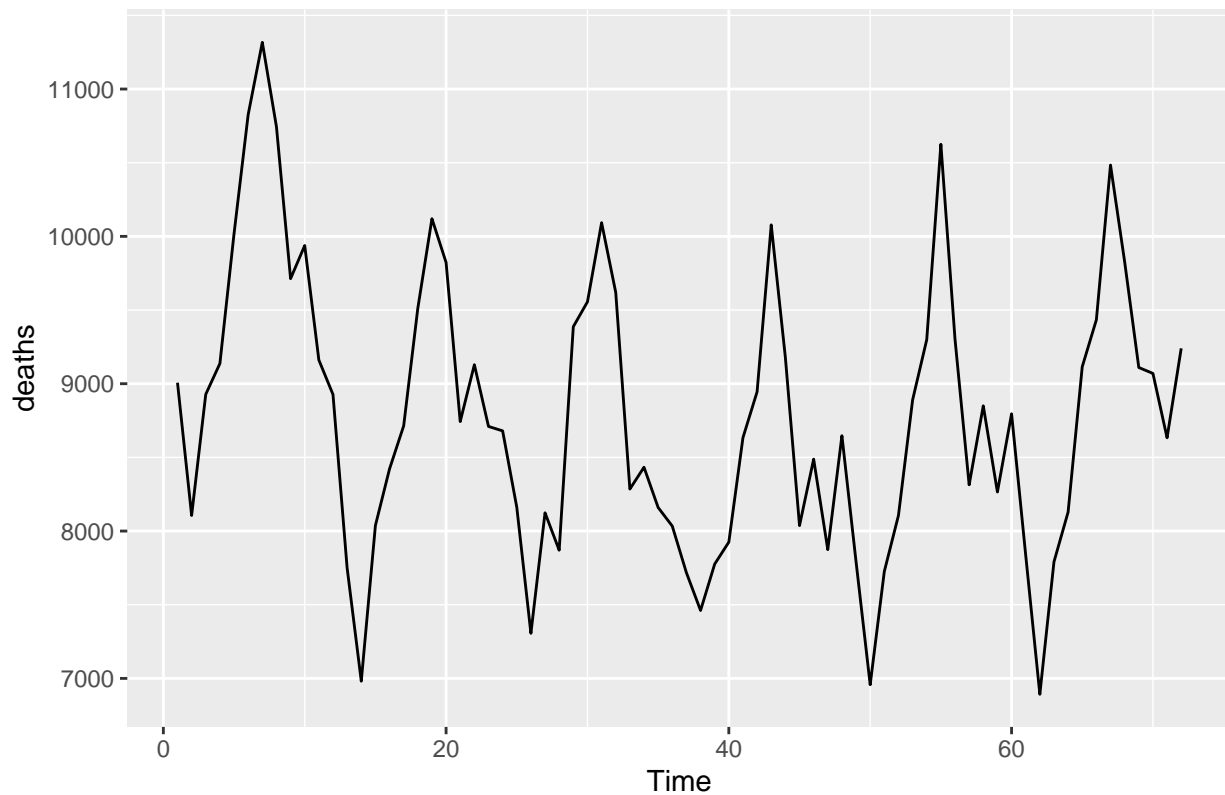
Start = 1

End = 72

Frequency = 1

```
[1] 9007 8106 8928 9137 10017 10826 11317 10744 9713 9938 9161 8927
[13] 7750 6981 8038 8422 8714 9512 10120 9823 8743 9129 8710 8680
[25] 8162 7306 8124 7870 9387 9556 10093 9620 8285 8433 8160 8034
[37] 7717 7461 7776 7925 8634 8945 10078 9179 8037 8488 7874 8647
[49] 7792 6957 7726 8106 8890 9299 10625 9302 8314 8850 8265 8796
[61] 7836 6892 7791 8129 9115 9434 10484 9827 9110 9070 8633 9240
```

```
# Plot the data for visualization
autoplot(deaths)
```



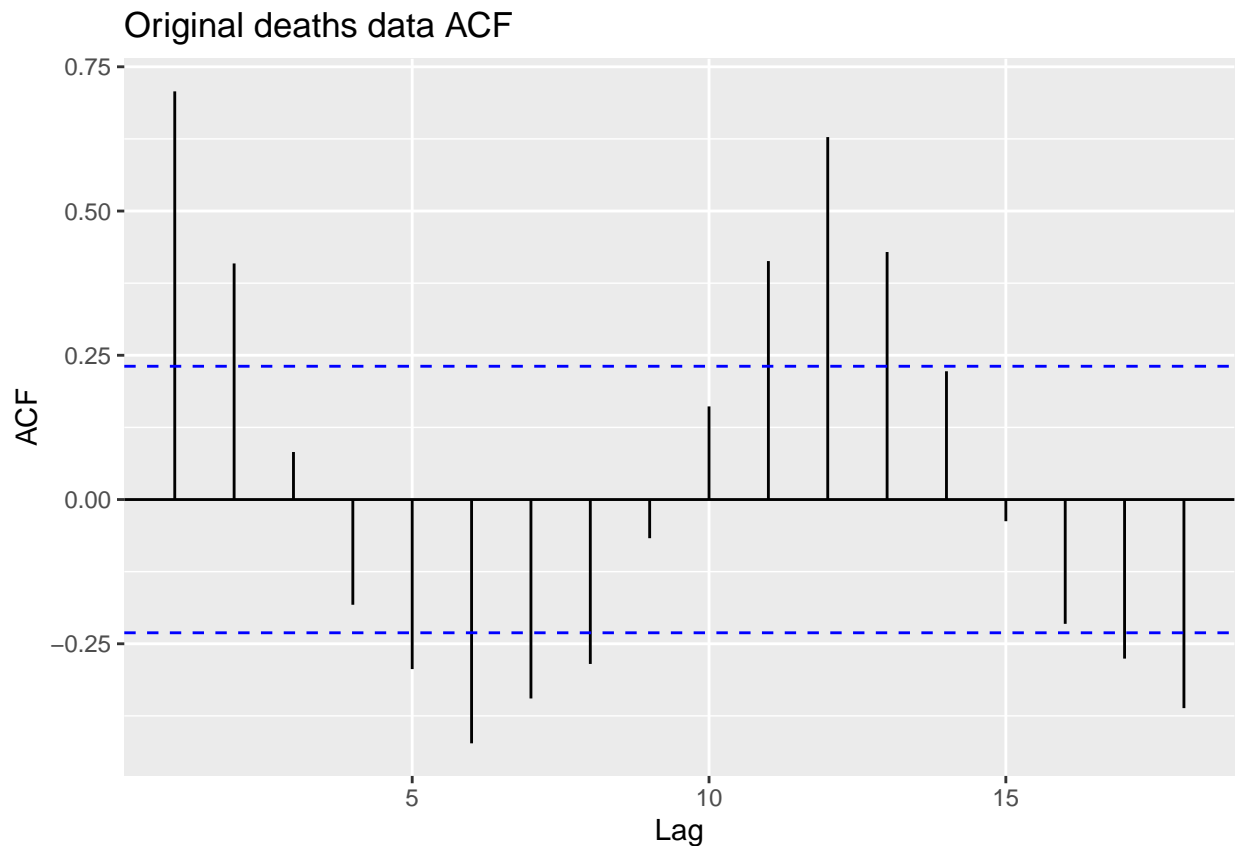
We can calculate the sample mean before differencing:

```
x_bar <- mean(deaths)
x_bar
```

```
[1] 8787.736
```

So $\hat{m}u = \bar{x}_{72} = 8787.736$.

```
# Original data ACF
deaths_acf <- ggAcf(deaths) + ggtitle("Original deaths data ACF")
deaths_acf
```



```
orig_sample_ACF_vals <- deaths_acf$data # extract the data object
orig_sample_ACF_vals <- c(1,orig_sample_ACF_vals$Freq) # extract autocorrelations
orig_sample_ACF_vals <- c(orig_sample_ACF_vals, orig_sample_ACF_vals[1:2]) # repeat some of these
orig_sample_ACF_vals
```

```
[1] 1.00000000 0.70739498 0.40914118 0.08232642 -0.18240432 -0.29389209
[7] -0.42256255 -0.34478440 -0.28495416 -0.06704366 0.16137269 0.41325445
[13] 0.62817012 0.42898283 0.22227129 -0.03760829 -0.21542724 -0.27571757
[19] -0.36154679 1.00000000 0.70739498
```

We will first difference to obtain the series of interest $\{\nabla\nabla_{12}X_t\}$:

```
deaths_diff_12 <- diff(deaths, 12) # difference of order 12
deaths_diff_1_12 <- diff(deaths_diff_12, 1) # diff again
deaths_diff_1_12 # Note we lost some data
```

Time Series:

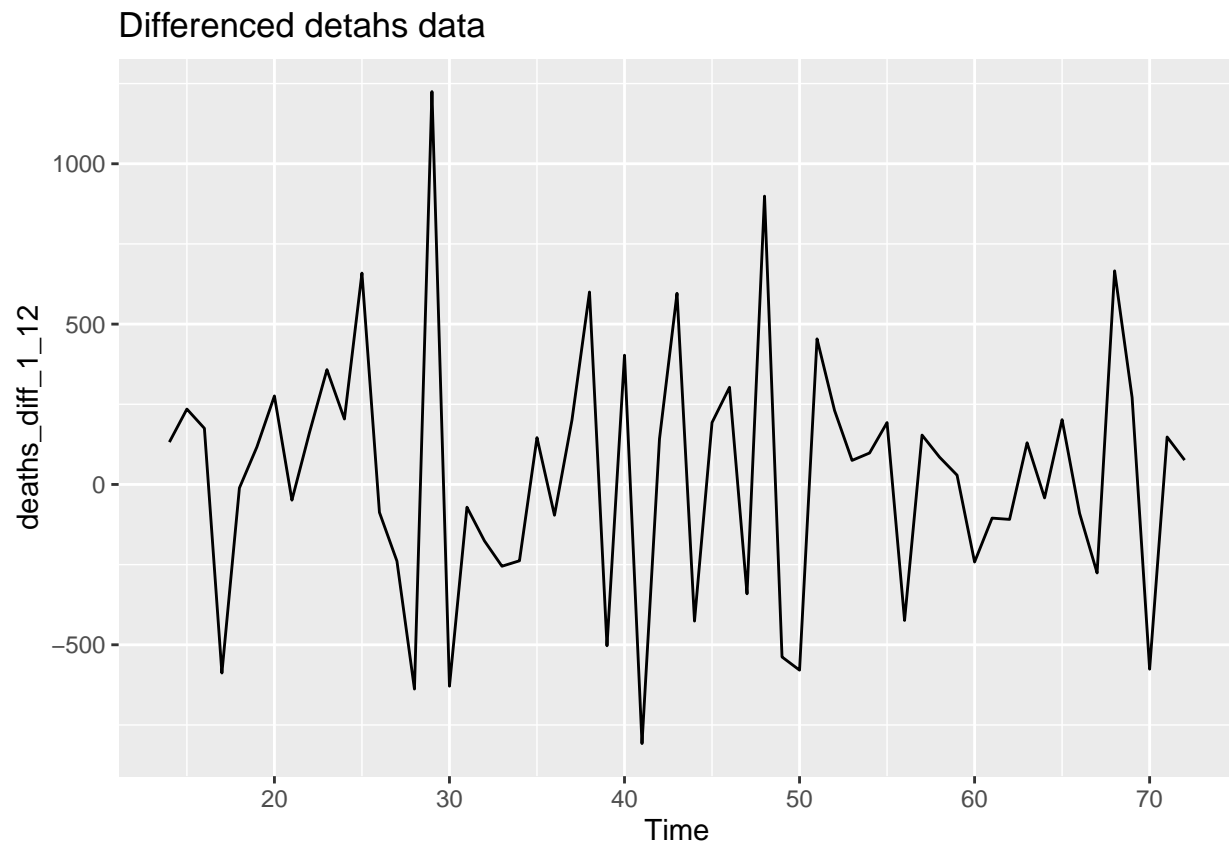
Start = 14

End = 72

Frequency = 1

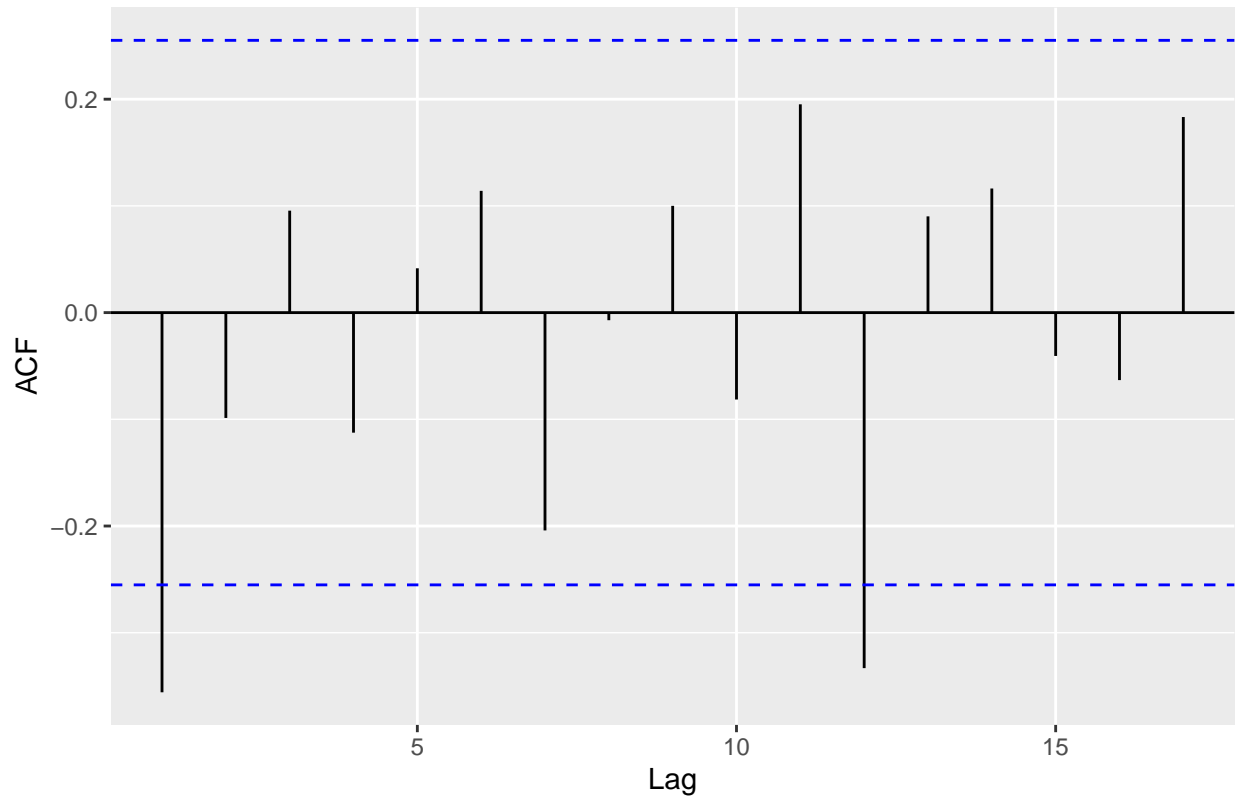
```
[1] 132 235 175 -588 -11 117 276 -49 161 358 204 659 -87 -239 -638
[16] 1225 -629 -71 -176 -255 -238 146 -96 201 600 -503 403 -808 142 596
[31] -426 193 303 -341 899 -538 -579 454 231 75 98 193 -424 154 85
[46] 29 -242 -105 -109 130 -42 202 -90 -276 666 271 -576 148 76
```

```
autoplot(deaths_diff_1_12) + ggtitle("Differenced deaths data")
```



```
deaths_diff_1_12_ACF = ggAcf(deaths_diff_1_12) + ggtitle("Differenced deaths data ACF")
deaths_diff_1_12_ACF
```

Differenced deaths data ACF



We can then extract sample autocovariances $\hat{\gamma}(h), 0 \leq h \leq 20$

```
sample_ACF_vals <- deaths_acf$data # extract the data object
sample_ACF_vals <- c(sample_ACF_vals$Freq) # extract autocorrelations
sample_ACF_vals <- c(sample_ACF_vals, sample_ACF_vals) # repeat some of these
sample_ACF_vals
```

```
[1] 0.70739498 0.40914118 0.08232642 -0.18240432 -0.29389209 -0.42256255
[7] -0.34478440 -0.28495416 -0.06704366 0.16137269 0.41325445 0.62817012
[13] 0.42898283 0.22227129 -0.03760829 -0.21542724 -0.27571757 -0.36154679
[19] 0.70739498 0.40914118 0.08232642 -0.18240432 -0.29389209 -0.42256255
[25] -0.34478440 -0.28495416 -0.06704366 0.16137269 0.41325445 0.62817012
[31] 0.42898283 0.22227129 -0.03760829 -0.21542724 -0.27571757 -0.36154679
```

Note that due to the differencing, we can only obtain $\hat{\gamma}(h), 0 \leq h \leq 18$, after which if the series is stationary, means that $\gamma(h) = \gamma(h, h + t)$, so that it simply gets repeated.

- (c) By equating $\hat{\gamma}(1), \hat{\gamma}(11)$ and $\hat{\gamma}(12)$ from part 3 with $\gamma(1), \gamma(2)$ and $\gamma(12)$ respectively from part (a), find a model of the form defined in (a) to represent $\{\nabla \nabla_{12} X_t\}$

- **NOTE:** Here I really wasn't sure which autocovariances I am supposed to take. From the assg, question 3.9, (a) we are asked to find the theoretical form of the autocovariances. Then, in (c) we are asked to compare these to the ACF of the **differenced** series, which doesn't make much sense? Instead, what makes more sense to me is to use the estimated ACF from the **original series** and compare these to the theoretical ACFs in part (a). Then, solve for θ_1, θ_{12} , and use these in the series $\nabla \nabla_{12} X_t$. So this is

what I've done here. If it's not what was supposed, then I should have used the values from the ACF of the latter series (again, doesn't make much sense to me). If that's the case, I should still get full points as this question was not clear.

So now we will extract the values of interest, $\hat{\gamma}(1)$, $\hat{\gamma}(11)$ and $\hat{\gamma}(12)$ from the original series.

```
gamma1 <- sample_ACF_vals[1]
gamma11 <- sample_ACF_vals[11]
gamma12 <- sample_ACF_vals[12]
gammas_hat <- c(gamma1,gamma11,gamma12)
names(gammas_hat) <- c("gamma1","gamma11","gamma12")
gammas_hat
```

```
      gamma1      gamma11      gamma12
0.7073950 0.4132545 0.6281701
```

Now, from (a), we have get the system of equations

$$\begin{cases} \sigma^2\theta_1 &= \hat{\gamma}_Y(1) \\ \sigma^2\theta_1\theta_{12} &= \hat{\gamma}_Y(11) \\ \sigma^2\theta_{12} &= \hat{\gamma}_Y(12) \end{cases}$$

So solving for the parameters we obtain

$$\begin{cases} \hat{\theta}_1 &= \frac{\hat{\gamma}_Y(1)}{\hat{\sigma}^2} = 0.6578705 \\ \hat{\theta}_{12} &= \frac{\hat{\gamma}_Y(12)}{\hat{\sigma}^2} = 0.5841921 \\ \hat{\sigma}^2 &= \frac{\hat{\gamma}_Y(1)\hat{\gamma}_Y(12)}{\hat{\gamma}_Y(11)} = 1.07528 \end{cases}$$

So the differenced series (after some calculations) has the form

$$\begin{aligned} \nabla_{12}X_{12} &= \hat{\theta}_1(Z_{t-1} - Z_{t-13}) + \hat{\theta}_{12}(Z_{t-12} - Z_{t-24}) \\ \nabla\nabla_{12}X_{12} &= \hat{\theta}_1(Z_{t-1} - Z_{t-2} - Z_{t-13} + Z_{t-14}) + \hat{\theta}_{12}(Z_{t-12} - Z_{t-13} - Z_{t-24} + Z_{t-25}) \end{aligned}$$

Final question

- (a) Remove the last 12 values from the **beer** data
- (b) Find an ARIMA model for the *logarithms* of the beerdata. Your analysis should include:
 - i) A logical explanation of the steps taken to choose the final model.
 - ii) Appropriate 95% bounds for the componenets of ϕ and θ .
 - iii) AN examination of the residuals to check for similarity of a white noise process.
 - iv) A graph of the series showing forecasts for the removed 12 values and 95% prediction bounds.
 - v) A table of the actual forecast errors, i.e. observed - predicted, for the removed 12 values
- (c) Repeat the steps in part (b), but instead use a classical decomposition approach by deseasonalizing, subtracting a quadratic trend, and then fitting and ARMA model to the reisudals. Then compare your forecast errors to those in part (b).

Preprocessing the data

We will first load the `beer` dataset

```
# Load the data into a ts object
beer_dat = dget("beer.Rput")
str(beer_dat)
```

```
Time-Series [1:422] from 1956 to 1991: 93.2 96 95.2 77.1 70.9 ...
```

Now we will consider the **logarithm** transformation of the data, which will be the focus of the analysis, then split the data into training and testing datasets:

```
# Preprocess data
beer_dat <- log(beer_dat) # logarithm transformation
beer_dat_train <- window(beer_dat, end=1990.1) # train dataset
beer_dat_test <- window(beer_dat, start=1990.1) # test dataset
str(beer_dat_train)
```

```
Time-Series [1:410] from 1956 to 1990: 4.53 4.56 4.56 4.35 4.26 ...
```

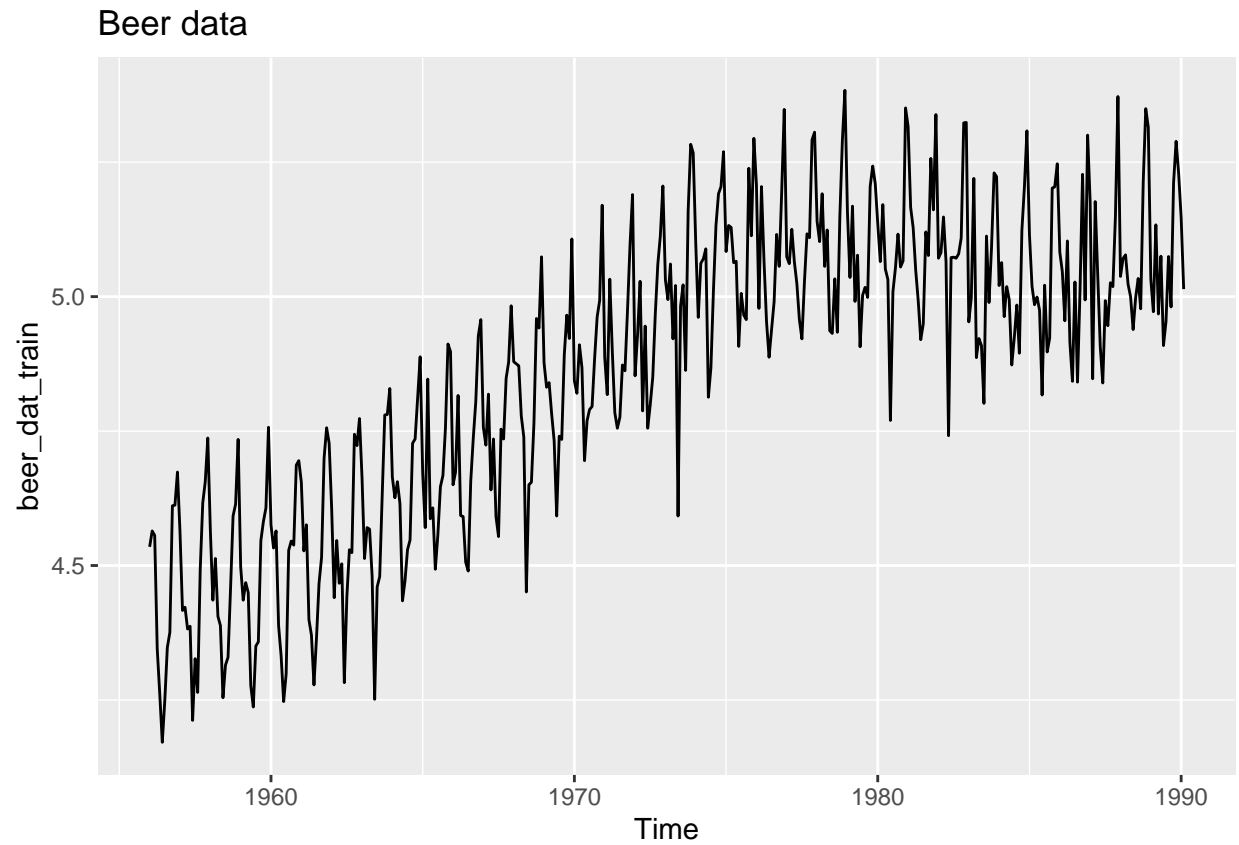
```
str(beer_dat_test)
```

```
Time-Series [1:12] from 1990 to 1991: 5.09 5.03 5.03 4.91 5 ...
```

Inspecting the data

Let's first take a look at the data

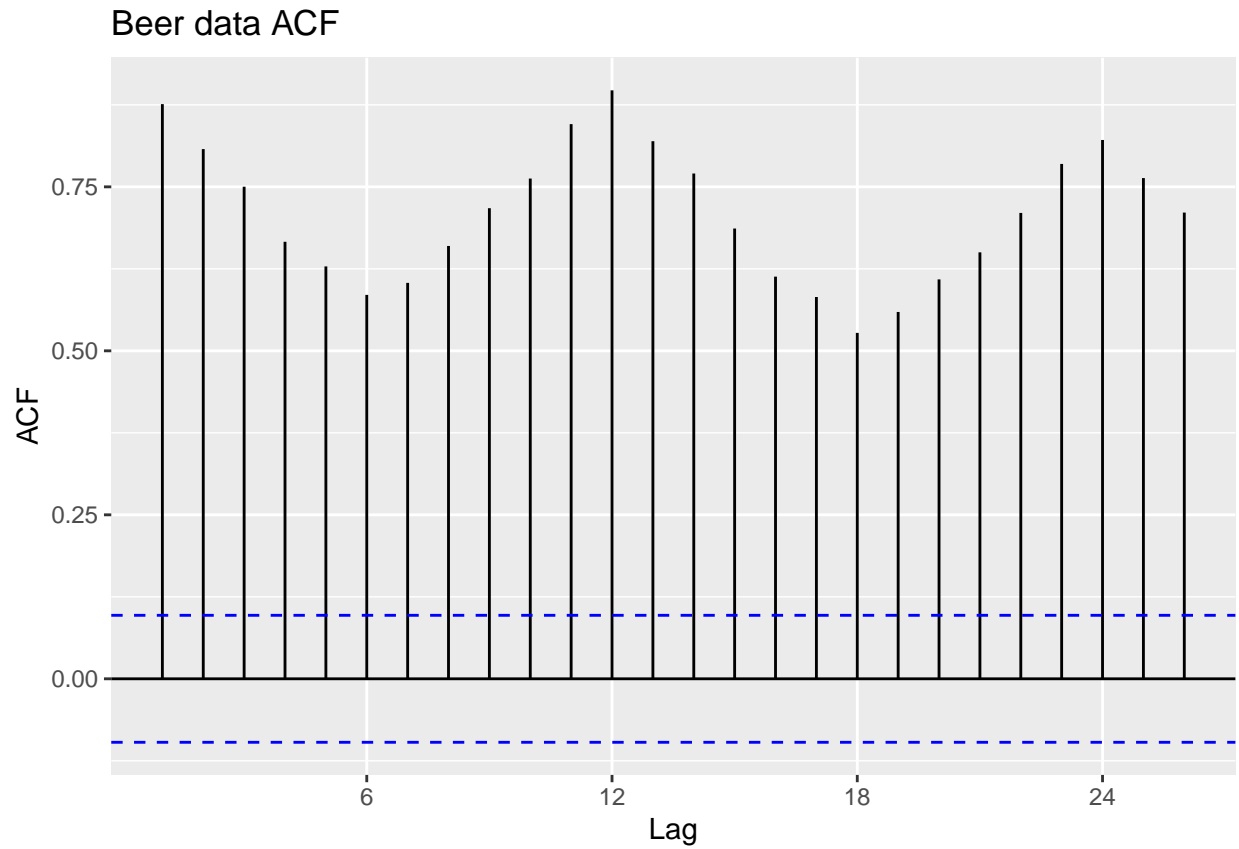
```
# Plot series
autoplot(beer_dat_train) + ggtitle("Beer data")
```



The data definitely does not look stationary. and some periodicity is also visible.

Let's take a look at the autocorrelation function:

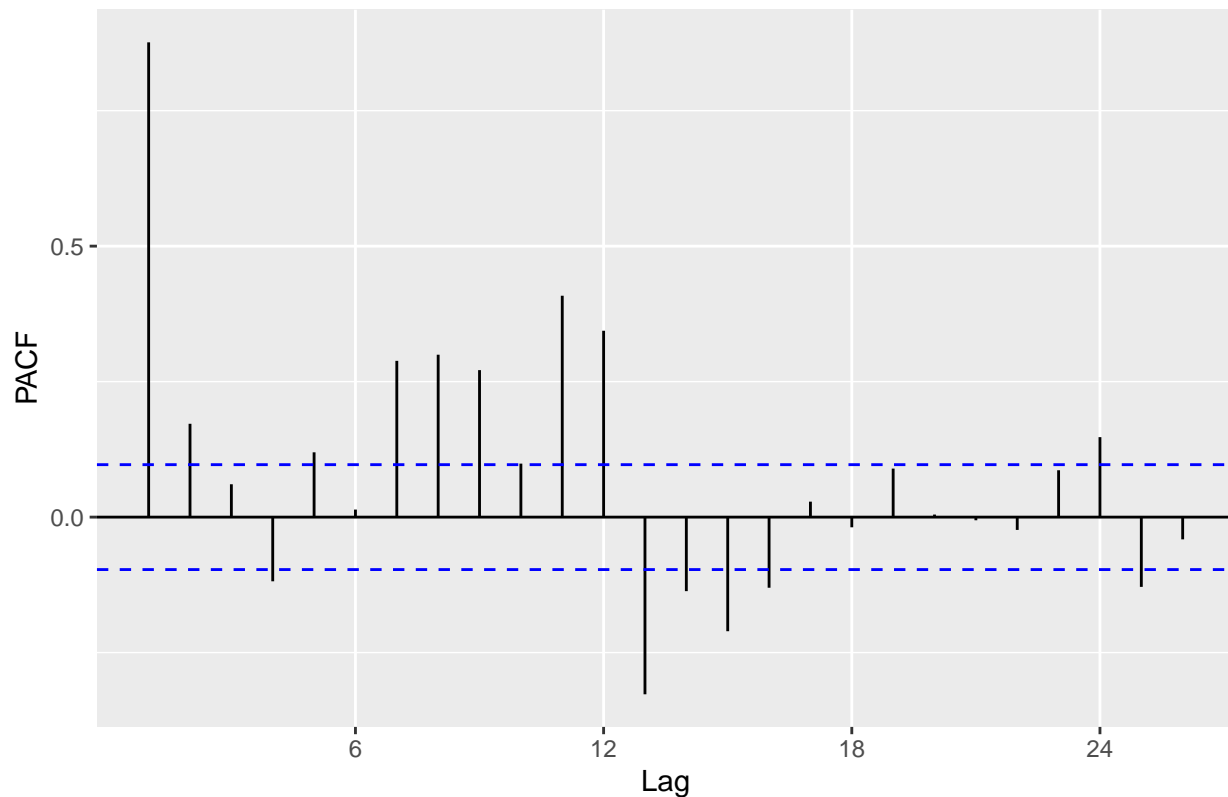
```
# Plot ACF  
ggAcf(beer_dat_train) + ggtitle("Beer data ACF")
```

Indeed, points at all lags are terribly correlated, and seasonality can be seen. This clearly tells this is definitely not a white noise $WN(0, \sigma^2)$ process.

```
# Plot PACF
ggPacf(beer_dat_train) + ggtitle("Beer data PACF")
```

Beer data PACF



However from the PACF plot, we observe that the correlation seems to decrease at further lags.

Fitting the ARIMA model

- Let $d \in \mathbb{Z}$. $\{X_t\}$ is said to be $\text{ARIMA}(p,d,q)$ iff $Y_t := (1 - B)^d X_t$ is a causal $\text{ARMA}(p,q)$ process.

We will now try to fit the best ARIMA model we can find to the data, using the `arma.sim` function (i.e. no ic selected)

```
# Fit an auto-arima model
beerdat_arma = auto.arma(beer_dat_train,
    seasonal=FALSE,
    stepwise=FALSE,
    ic = c("aicc", "aic", "bic") ,
    approximation=FALSE,
    trace=TRUE)
```

ARIMA(0,1,0)		: -472.2122
ARIMA(0,1,0)	with drift	: -470.223
ARIMA(0,1,1)		: -495.6891
ARIMA(0,1,1)	with drift	: -493.7302
ARIMA(0,1,2)		: -534.0083
ARIMA(0,1,2)	with drift	: -539.507
ARIMA(0,1,3)		: -544.2853

ARIMA(0,1,3)	with drift	: -549.5625
ARIMA(0,1,4)		: -570.0043
ARIMA(0,1,4)	with drift	: -574.6963
ARIMA(0,1,5)		: -569.6033
ARIMA(0,1,5)	with drift	: -573.9958
ARIMA(1,1,0)		: -492.1827
ARIMA(1,1,0)	with drift	: -490.2058
ARIMA(1,1,1)		: -555.8088
ARIMA(1,1,1)	with drift	: -560.3982
ARIMA(1,1,2)		: -553.7811
ARIMA(1,1,2)	with drift	: -558.3878
ARIMA(1,1,3)		: -554.0459
ARIMA(1,1,3)	with drift	: -558.6354
ARIMA(1,1,4)		: -570.0013
ARIMA(1,1,4)	with drift	: -574.3103
ARIMA(2,1,0)		: -494.7938
ARIMA(2,1,0)	with drift	: -492.8235
ARIMA(2,1,1)		: -503.6188
ARIMA(2,1,1)	with drift	: -501.6311
ARIMA(2,1,2)		: -556.337
ARIMA(2,1,2)	with drift	: -561.0268
ARIMA(2,1,3)		: -554.9761
ARIMA(2,1,3)	with drift	: -559.5321
ARIMA(3,1,0)		: -494.544
ARIMA(3,1,0)	with drift	: -492.5492
ARIMA(3,1,1)		: -502.1042
ARIMA(3,1,1)	with drift	: -500.1003
ARIMA(3,1,2)		: -555.9742
ARIMA(3,1,2)	with drift	: -560.3795
ARIMA(4,1,0)		: -504.1889
ARIMA(4,1,0)	with drift	: -502.2292
ARIMA(4,1,1)		: -603.5335
ARIMA(4,1,1)	with drift	: -610.4952
ARIMA(5,1,0)		: -503.2939
ARIMA(5,1,0)	with drift	: -501.3409

Best model: ARIMA(4,1,1) with drift

The fitting method employed is conditional-sum-of-squares to find starting values, then maximum likelihood. In this case, as the `ic` (information criterion) was not specified, the algorithm performed model selection according to all of the three criteria: AIC, AICC and BIC. The values displayed on the right-most column are $-2\ell(x; \hat{\phi}, \hat{\theta}, \hat{\sigma}^2)$ for fitting the different models. Whenever we obtained `Inf`. It means that the model was so terrible that it could not find appropriate initial values, and it diverged.

Inspecting the chosen ARIMA(4,1,1) model

Now let's inspect our model:

```
beerdat_arima_411 <- beerdat_arima
beerdat_arima_411
```

Series: beer_dat_train
ARIMA(4,1,1) with drift

Coefficients:

	ar1	ar2	ar3	ar4	ma1	drift
	0.4655	0.0537	0.0512	-0.3486	-0.9360	0.0017
s.e.	0.0470	0.0517	0.0516	0.0469	0.0126	0.0005

sigma² estimated as 0.01282: log likelihood=312.39
AIC=-610.77 AICc=-610.5 BIC=-582.68

This says that we have a lag 1 difference giving an ARMA(3,1) model. From this, we see all the estimated parameters $\hat{\phi}_1, \hat{\phi}_2, \hat{\phi}_3, \hat{\phi}_4, \hat{\theta}_1$ and $\hat{\sigma}^2$. (Note how little variance there is!)

Parameters

We can construct 95% confidence intervals for these:

```
## Calculate the wald confidence interval
params_arima <- beerdat_arima_411$coef # extract parameters
se_arima <- sqrt(diag(beerdat_arima_411$var.coef)) # extract s.e. for each
z <- qnorm(0.975) # this is the alpha/2 quantile
c.upper <- params_arima + z*se_arima # upper bound
c.lower <- params_arima - z*se_arima # lower bound
CI_arima_411 <- cbind(params_arima, se_arima, c.lower, c.upper)
colnames(CI_arima_411) <- c("parameter", "s.d.", "2.5 %", "97.5 %")
CI_arima_411
```

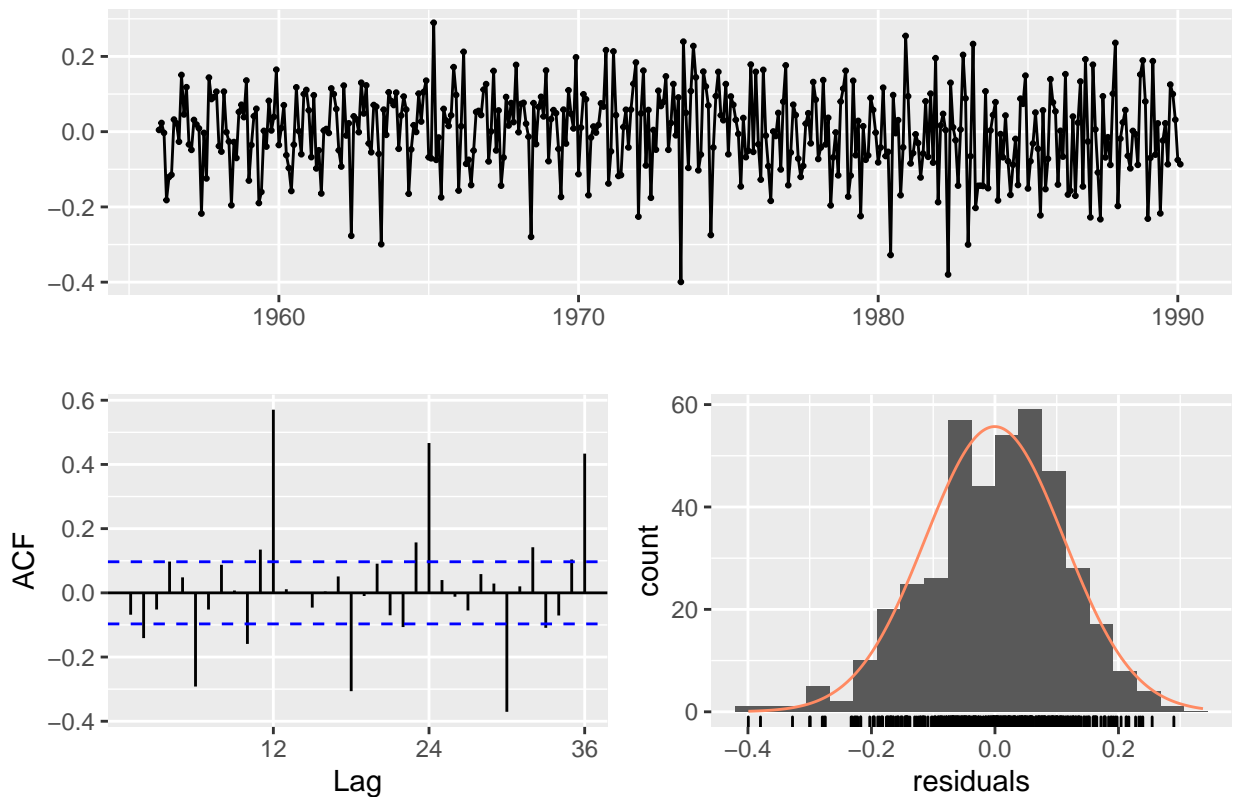
	parameter	s.d.	2.5 %	97.5 %
ar1	0.465458908	0.0470214978	0.3732984661	0.557619350
ar2	0.053715988	0.0517065296	-0.0476269473	0.155058924
ar3	0.051233576	0.0516342440	-0.0499676828	0.152434834
ar4	-0.348570052	0.0468793510	-0.4404518912	-0.256688212
ma1	-0.935950001	0.0126007088	-0.9606469365	-0.911253066
drift	0.001658102	0.0004784955	0.0007202684	0.002595936

Residuals inspection and diagnostics

Let's observe the residuals from the resulting model:

```
checkresiduals(beerdat_arima_411)
```

Residuals from ARIMA(4,1,1) with drift



Ljung-Box test

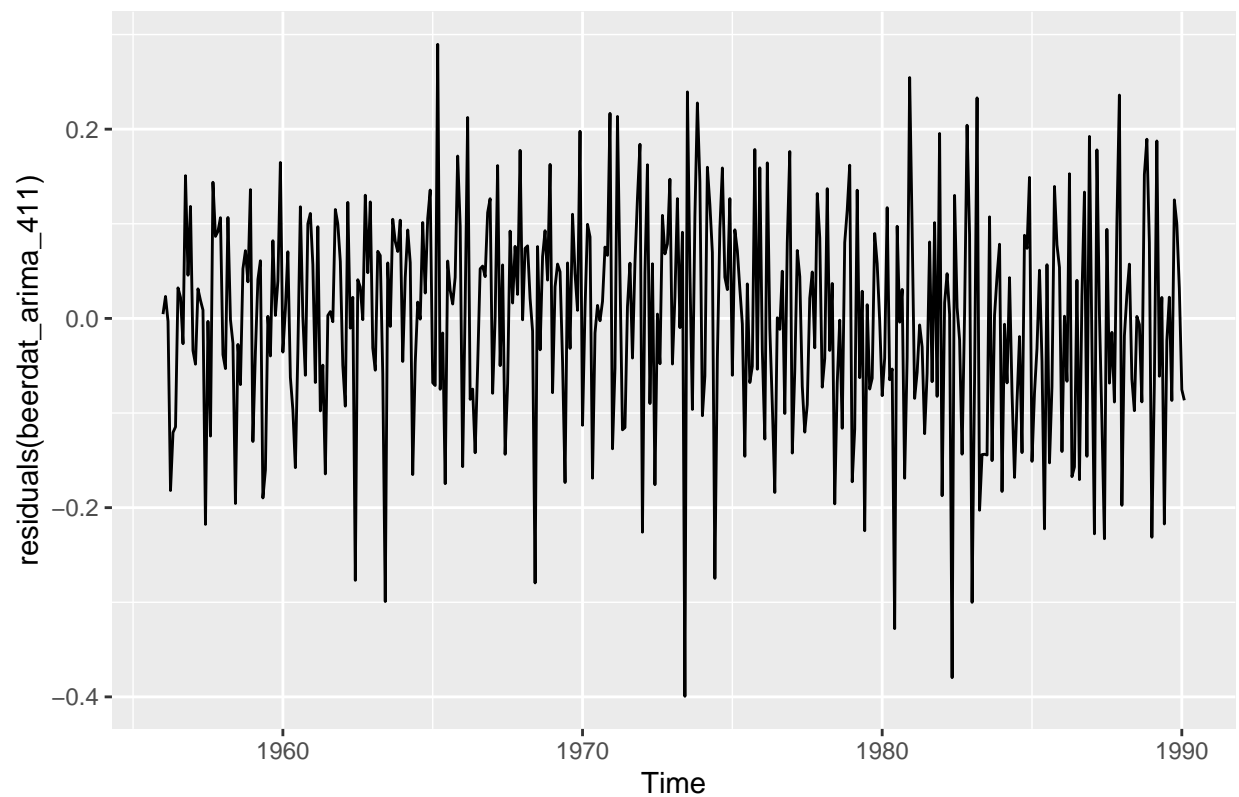
data: Residuals from ARIMA(4,1,1) with drift

Q* = 371.97, df = 18, p-value < 2.2e-16

Model df: 6. Total lags used: 24

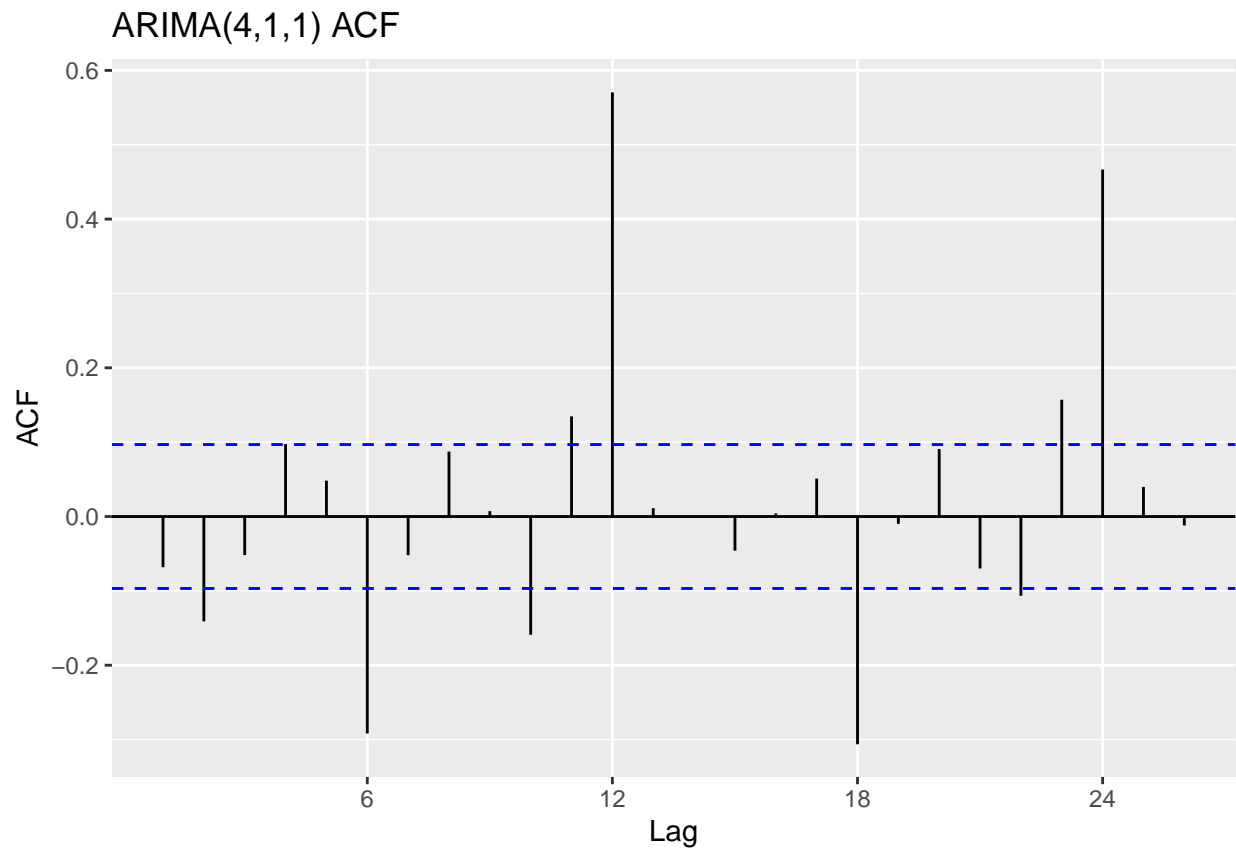
At first sight, it would seem that the residuals look somehow stationary and zero centered, as well as approximatedly normal. However, it would seem that there is still seasonality floating around! Indeed, from the Ljung-Box test, we see that we strongly reject the hypothesis of stationarity. **HOWEVER!** This is the wrong approach for testing stationarity for an ARIMA(p,d,q) model, so we wil come back to this.

```
# Plot the resulting model
autoplot(residuals(beerdat_arima_411))
```

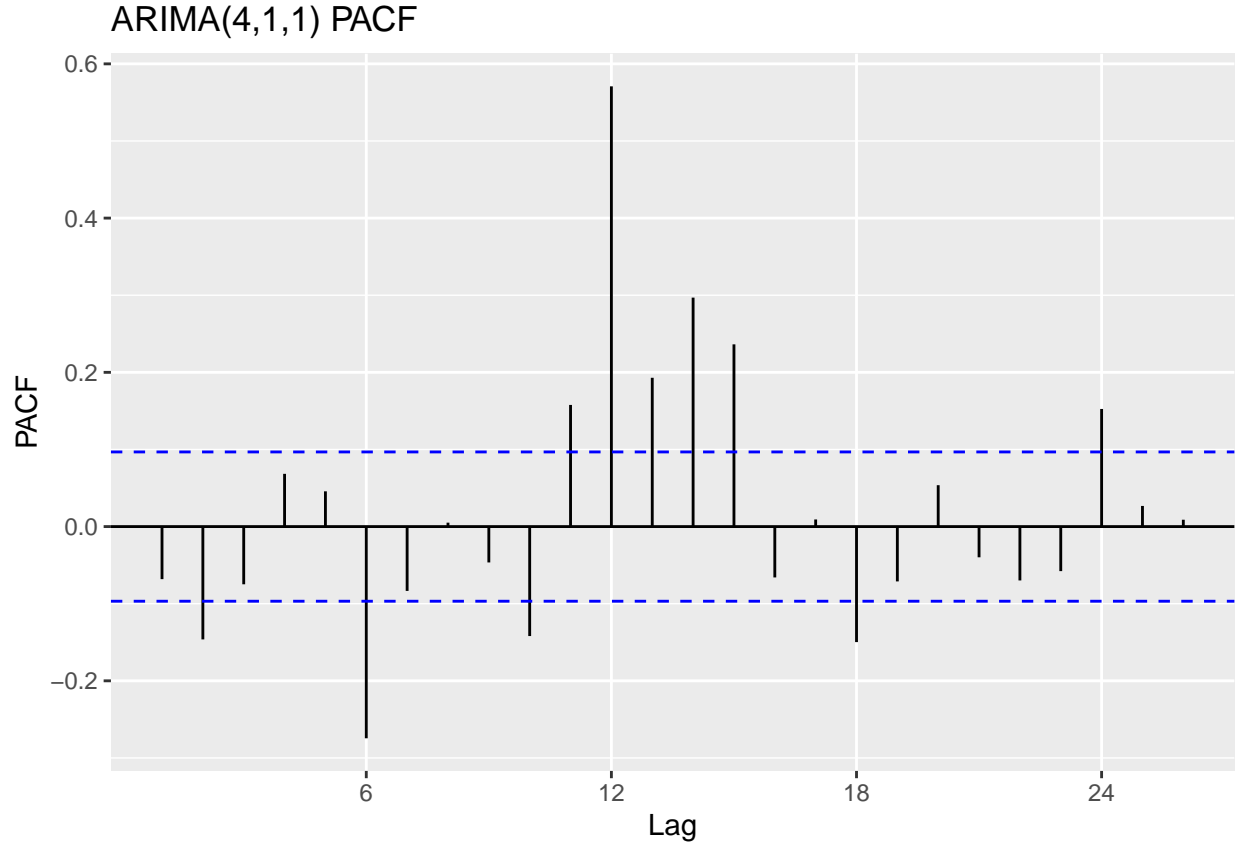


We observe that indeed the model seems to be a lot more centered towards 0, and seems a little bit more like white noise. Let us check this with the ACF and PACF plots:

```
# Plot ACF  
ggAcf(residuals(beerdatt_arima_411)) + ggtitle("ARIMA(4,1,1) ACF")
```



```
# Plot PACF  
ggPacf(residuals(beerdat_arima_411)) + ggtitle("ARIMA(4,1,1) PACF")
```



Inspecting the roots

Indeed, there is too much of a seasonal component floating around.

All of this seems to indicate the model is not appropriate, however, recall the definitions of stationarity, invertibility and causality for an ARMA(p,q) process:

- **Stationarity:** A solution to an ARMA(p,q) exists and is **stationary** iff

$$\Phi(z) = 1 - \sum_{j=1}^p \phi_j z^j \neq 0, \quad \forall |z| = 1 \text{ (for } z \text{ on the unit circle)}$$

i.e. we want **no roots** of $\Phi(z)$ on the unit circle.

- **Causality:** An ARMA(p,q) is **causal** if $\exists \{\psi_j\}$ such that

$$\sum_{j=1}^{\infty} |\psi_j| < \infty \quad \text{and} \quad X_t = \sum_{j=1}^{\infty} \psi_j Z_{t-j} \quad (\text{i.e. } X_t \text{ only depends on the past})$$

equivalently, iff

$$\Phi(z) = 1 - \sum_{j=1}^p \phi_j z^j \neq 0, \quad \forall |z| \leq 1 \text{ (for } z \text{ on and inside the unit circle)}$$

- **Invertibility:** An ARMA(p,q) process $\{X_t\}$ is invertible if \exists coefficients $\{\pi_j\}$ such that

$$\sum_{j=1}^{\infty} |\pi_j| < \infty \quad \text{and} \quad Z_t = \sum_{j=1}^{\infty} \pi_j X_{t-j}$$

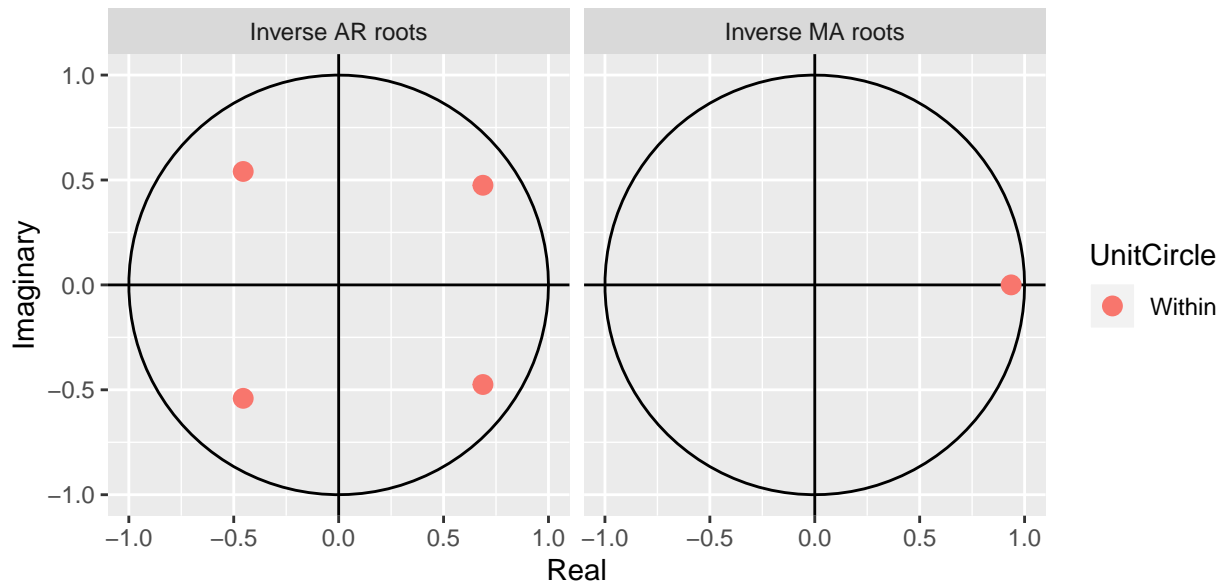
equivalently, we want

$$\Theta(z) = 1 + \sum_{j=1}^q \theta_k z^j \neq 0, \quad \forall |z| \leq 1$$

\equiv no roots of $\Theta(z)$ on or inside the unit circle.

So let's inspect the roots of what would be the AR(4) AND MA(1) polynomials:

```
# Plot the toods
autoplot(beerdat_arima_411)
```



We observe that all the inverse AR and MA roots lie inside the unit circle respectively, which indicates that our ARMA(4,1) process we would obtain by differencing once the ARMA(4,1,1) is a stationary, causal and invertible process.

Testing the ARIMA(4,1,1) stationarity

Since we **cannot** use likelihood-based approaches for the ARIMA, instead, we will perform ADP and KPSS tests before, and after taking one difference of the ARMA(4,1,1) process. Recall the respective hypotheses

for each of the tests:

- **ADP**

$$\begin{cases} H_0 : \text{The process is non-stationary (i.e. at least one unit root)} \\ H_a : \text{The process is stationary} \end{cases}$$

- **KSPP**

$$\begin{cases} H_0 : \text{The process is stationary (i.e. no unit roots)} \\ H_a : \text{The process is non-stationary} \end{cases}$$

```
# Test with a bunch of different k's ? (bigger augmented versions)
adf.test(residuals(beerdat_arima_411))
```

```
Warning in adf.test(residuals(beerdat_arima_411)): p-value smaller than printed
p-value
```

Augmented Dickey-Fuller Test

```
data: residuals(beerdat_arima_411)
Dickey-Fuller = -10.12, Lag order = 7, p-value = 0.01
alternative hypothesis: stationary
```

```
kpss.test(residuals(beerdat_arima_411))
```

```
Warning in kpss.test(residuals(beerdat_arima_411)): p-value smaller than printed
p-value
```

KPSS Test for Level Stationarity

```
data: residuals(beerdat_arima_411)
KPSS Level = 0.74841, Truncation lag parameter = 5, p-value = 0.01
```

So in this case the ADF test tells us that the process is stationary, while the second one tells us it is not. Therefore we should try differencing to see whether this solves the issue.

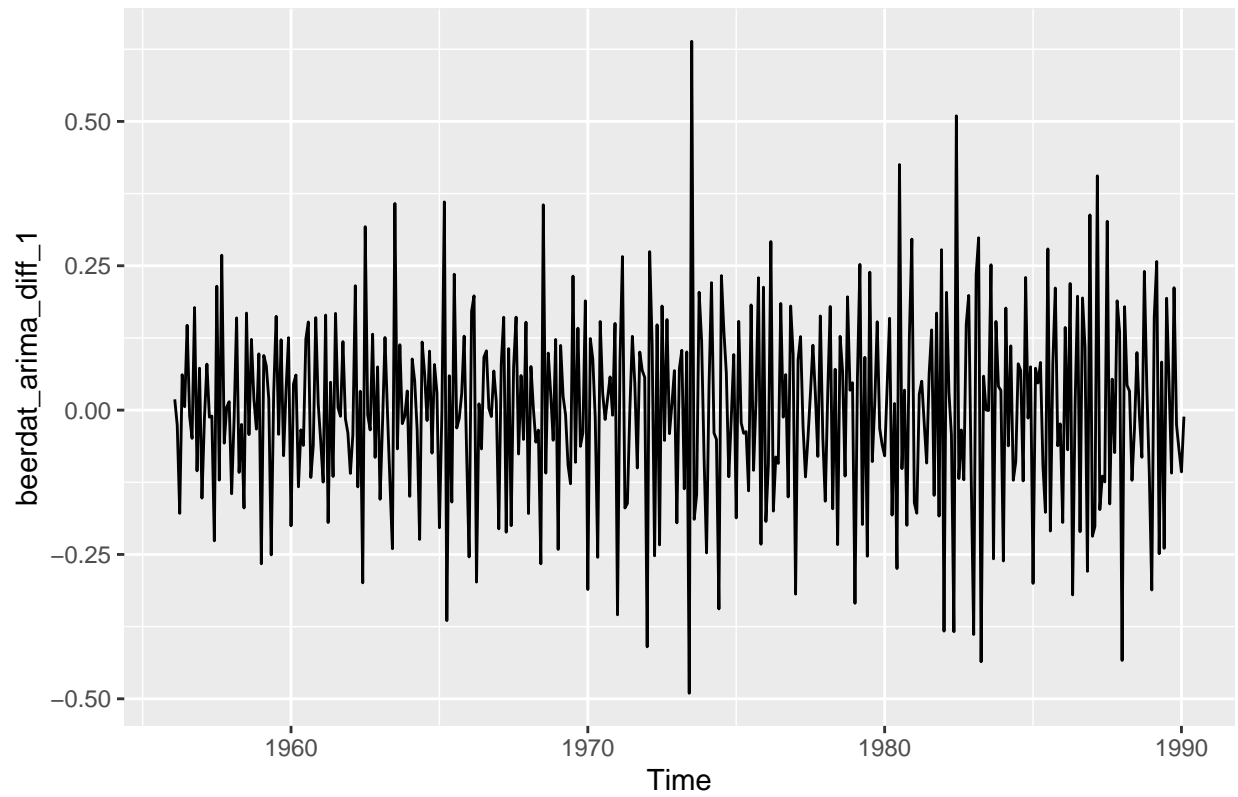
Differencing the ARIMA(4,1,1)

The new process is

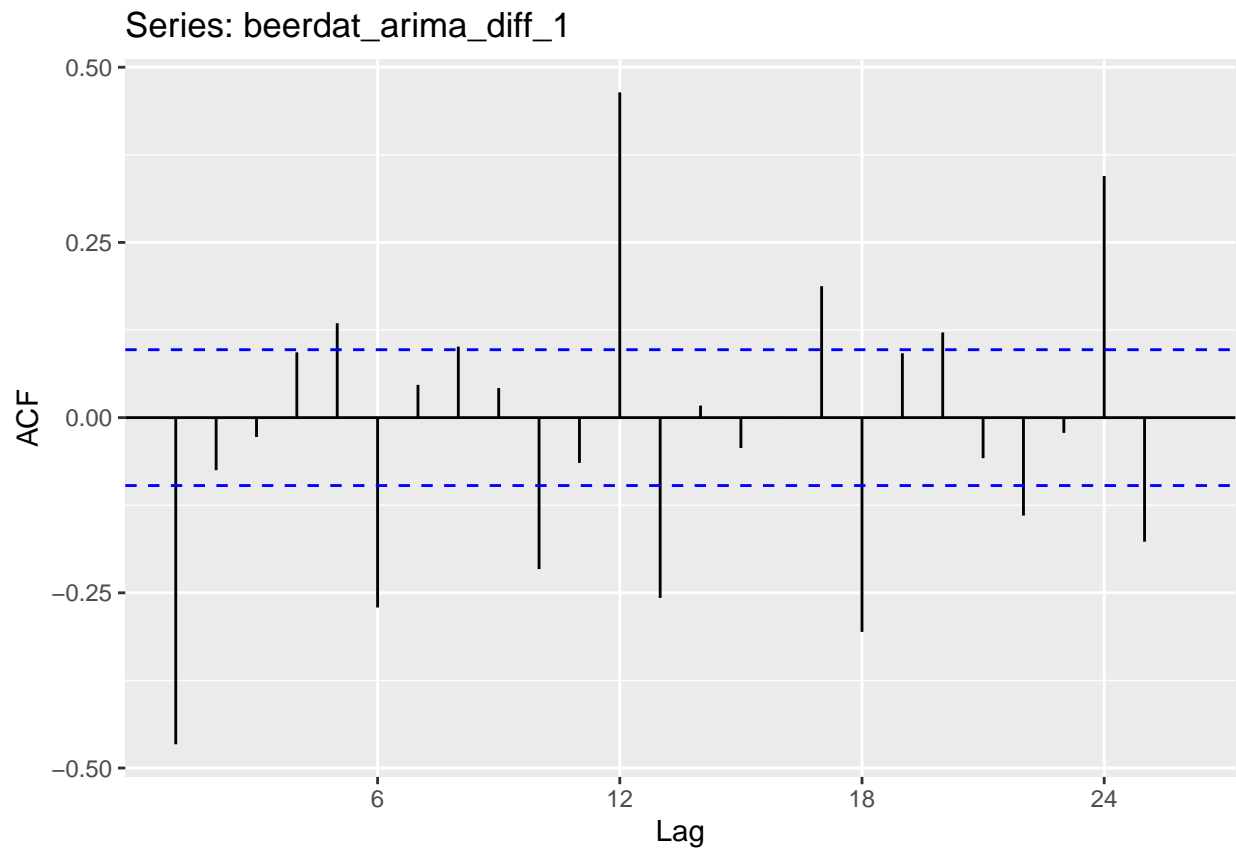
```
# differentiate the residuals form before
beerdat_arima_diff_1 <- diff(residuals(beerdat_arima_411), 1)
```

Let's inspect the plot, ACF and PACF

```
# Plot the series  
autoplot(beerdat_arima_diff_1)
```

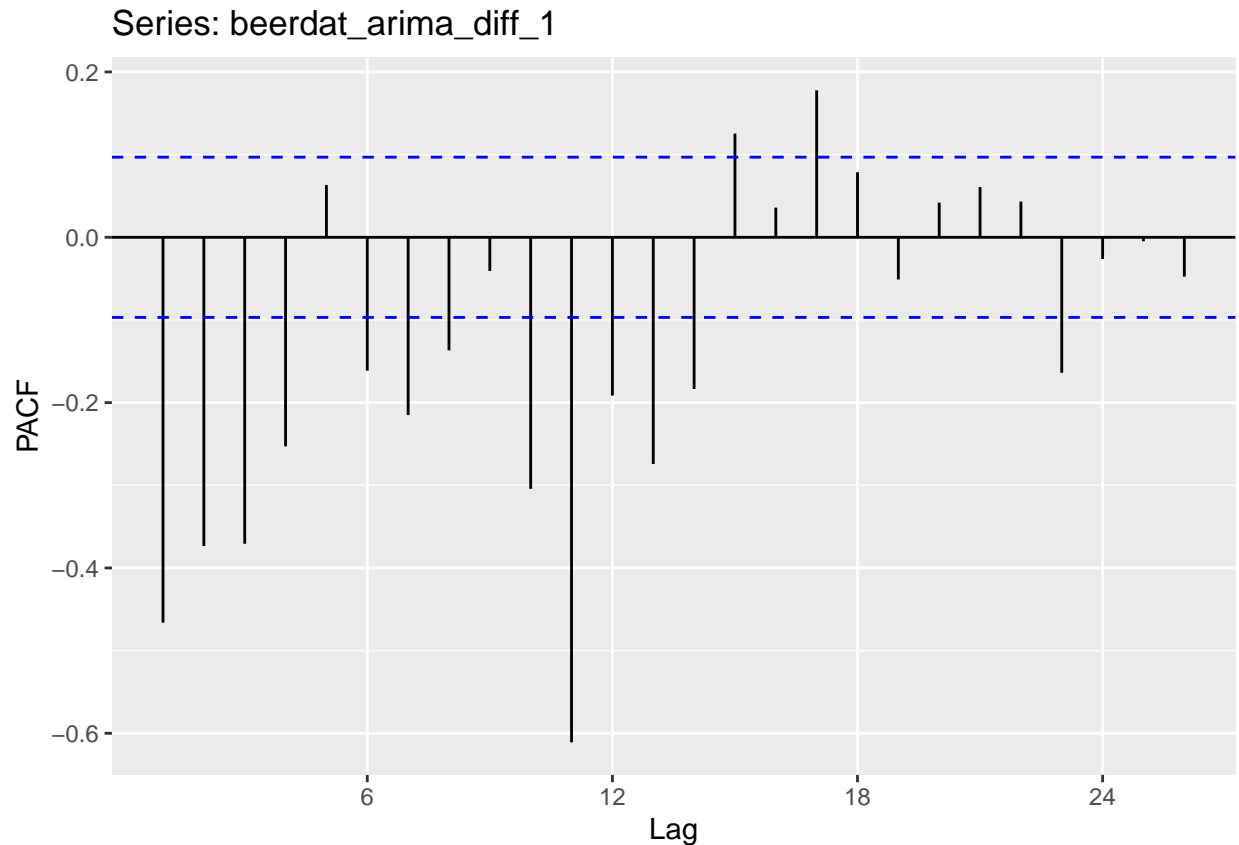


```
# Plot ACF  
ggAcf(beerdat_arima_diff_1)
```



This is still not good! :(

```
# Plot ACF  
ggPacf(beerdat_arima_diff_1)
```



In fact it looks like we made things worse. That PACF doesn't look healthy at all. This model is not good.

Attempting to remove seasonality from the ARIMA(1,1,1)

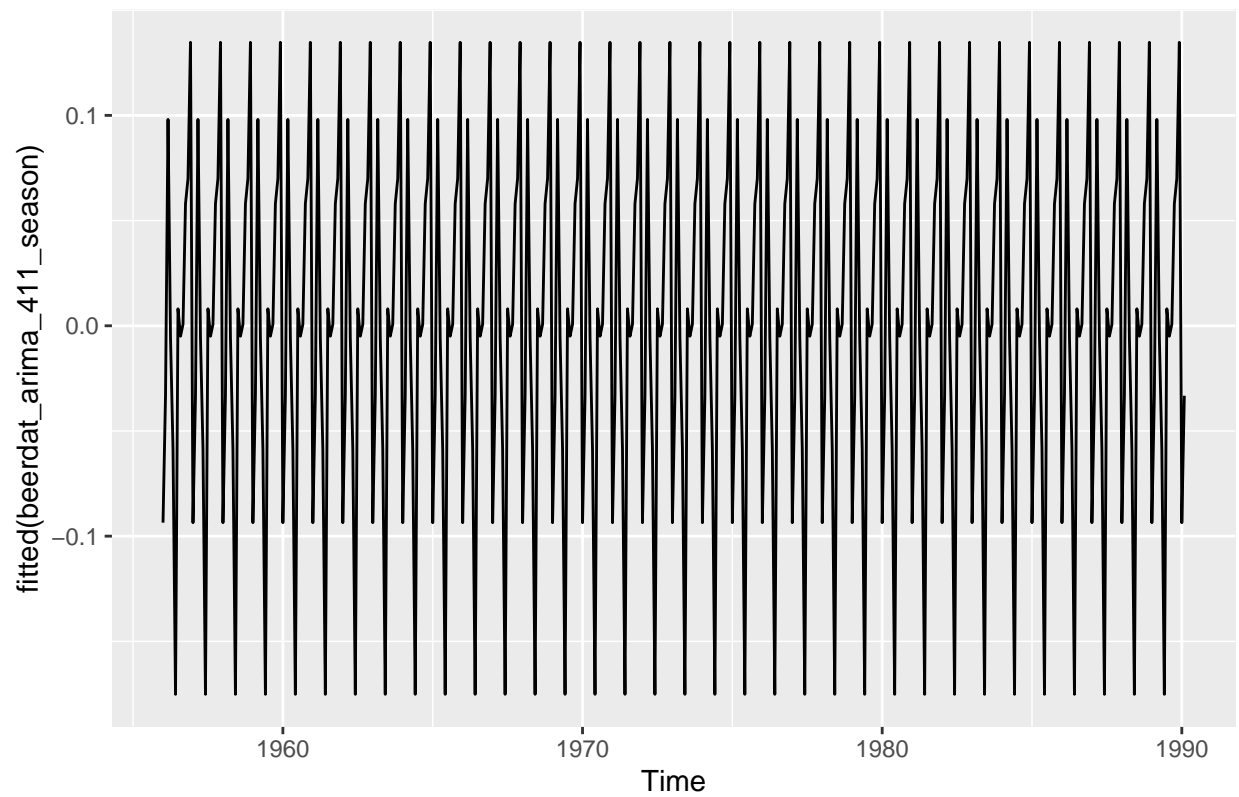
Let's try removing seasonality from the ARIMA(4,1,1) **before** differencing to see if this helps.

```
# frequency obtains the lag?
frequency(residuals(beerdat_arima_411))
```

```
[1] 12
```

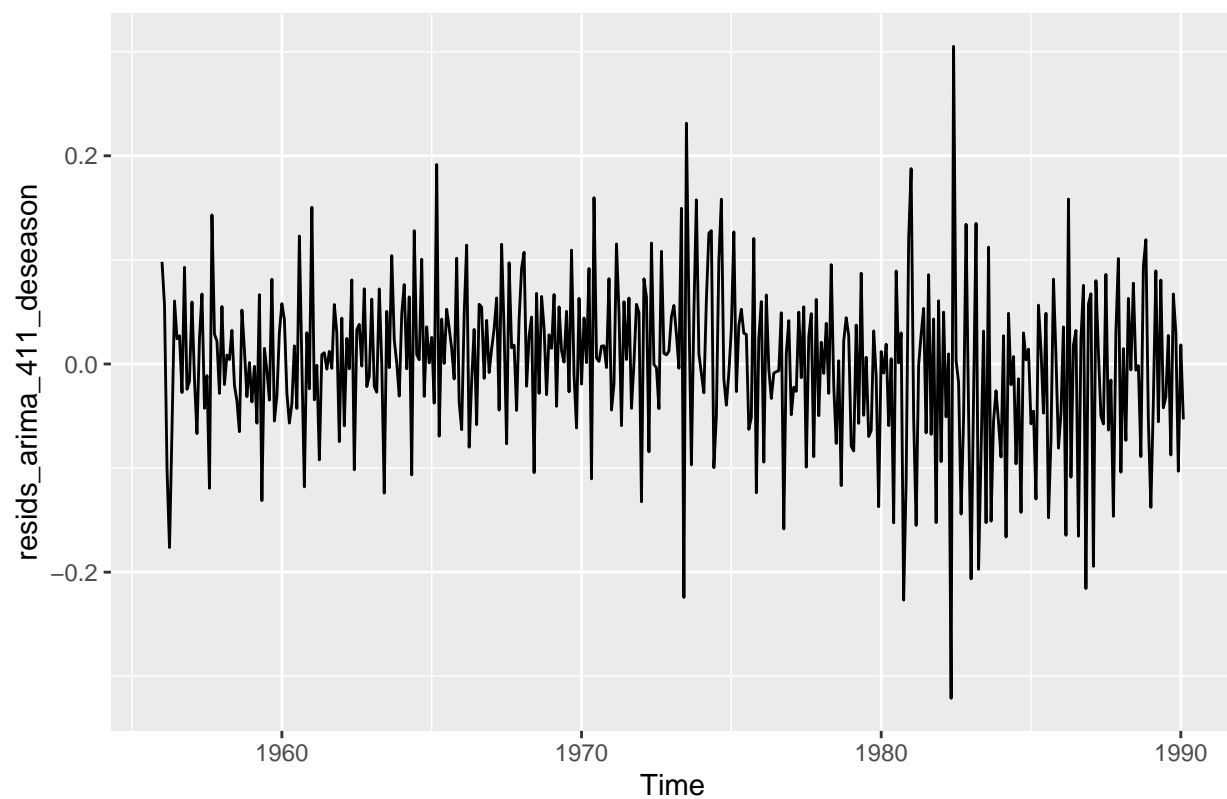
This is an indication that we have a lag every 12 months (indeed, every year!)

```
# attempt to calculate the de-seasoned data
beerdat_arima_411_season <- tslm(residuals(beerdat_arima_411) ~ season)
autoplot(fitted(beerdat_arima_411_season))
```



Now let's remove this from the original series

```
# Remove seasonal component  
resids_arima_411_deseason <- ( residuals(beerdat_arima_411) - fitted(beerdat_arima_411_season) )  
autoplot(resids_arima_411_deseason)
```

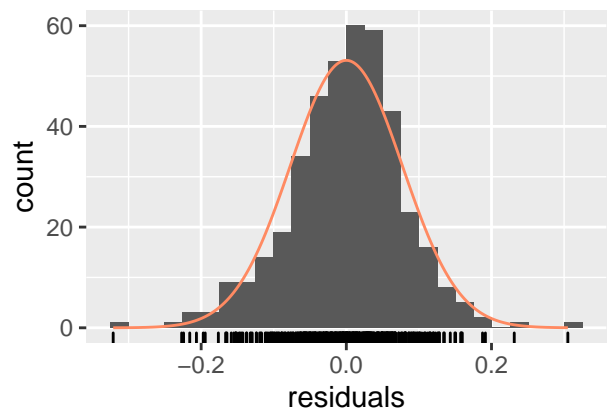
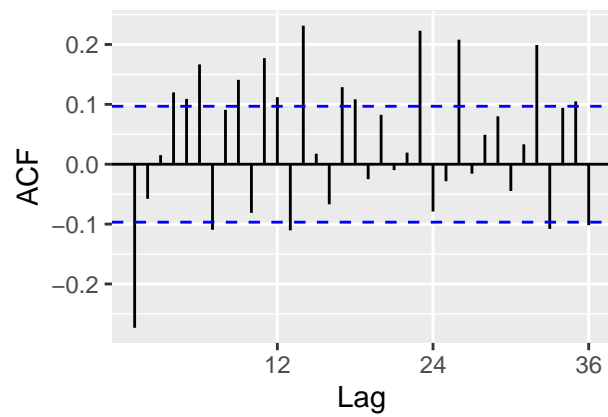
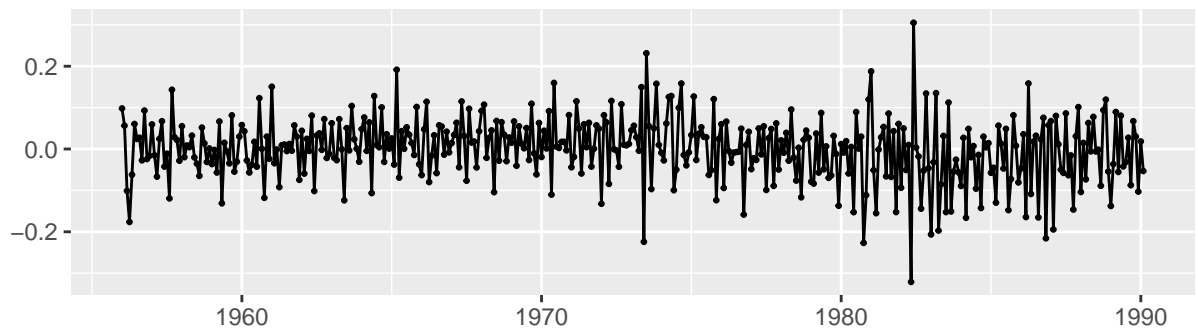


Now let's inspect the residuals

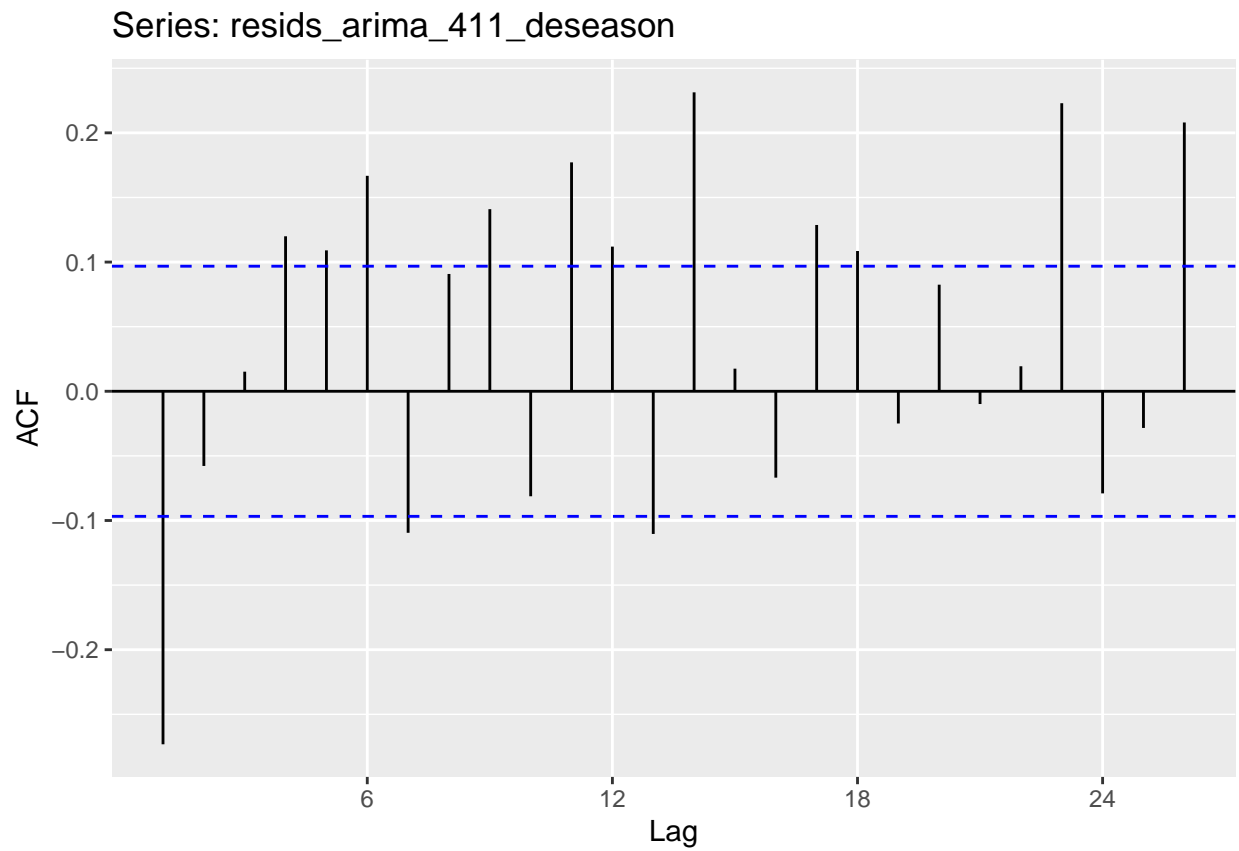
```
checkresiduals(resids_arima_411_deseason)
```

```
Warning in modeldf.default(object): Could not find appropriate degrees of  
freedom for this model.
```

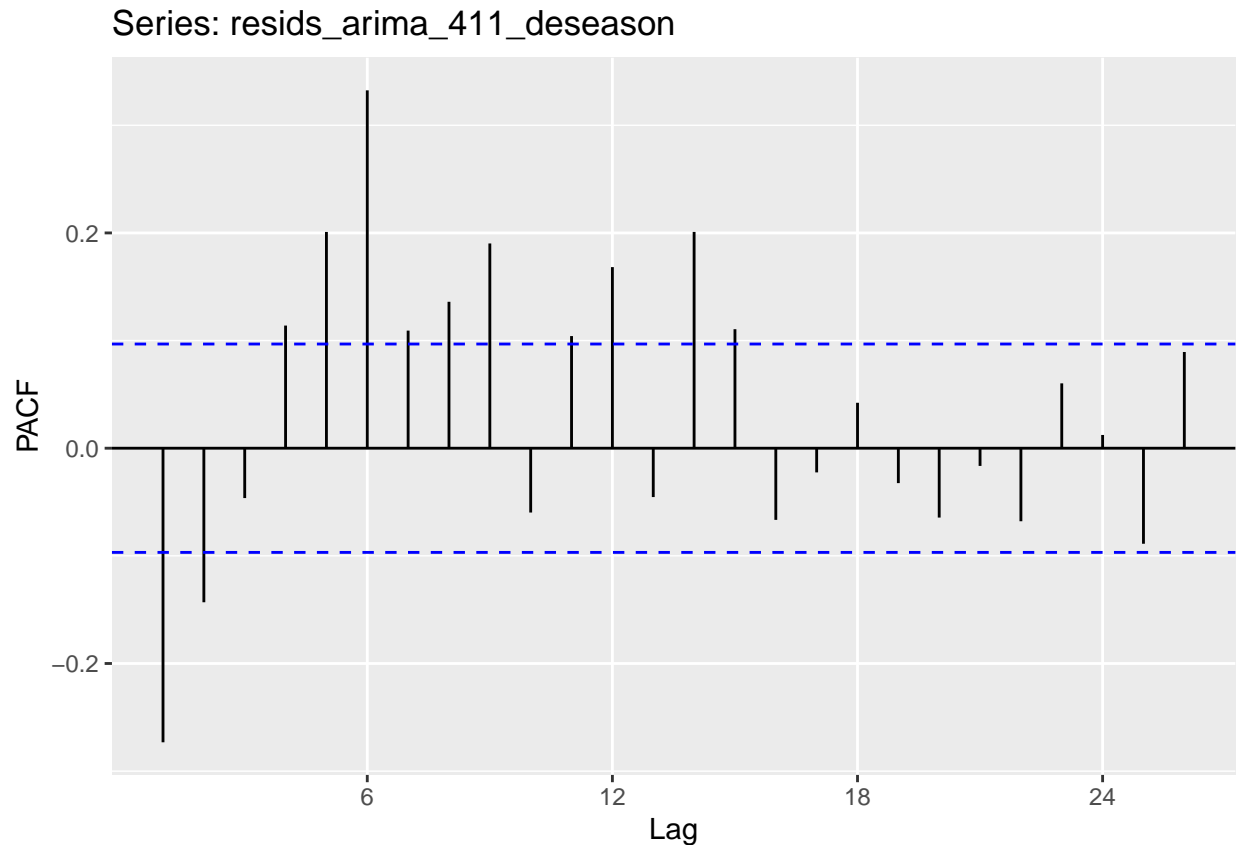
Residuals



```
# Plot ACF  
ggAcf(resids_arima_411_deseason)
```

```
# Plot PACF  
ggPacf(resids_arima_411_deseason)
```



Although we now have a series that looks slightly a bit more like white noise, the ACF plot still doesn't look that healthy. It appears that now a new trend has surged, which indicates that perhaps we should have removed it from the beginning instead. For now, we will simply see the forecasted values of the ARIMA(4,1,1) model.

Forecasting

Just for clarity, let's re-fit the ARIMA(4,1,1) model with the training data and predict on the test data:

```
# Fit the ARIMA model on training data
beerdat_arima_411 = auto.arima(beer_dat_train,
                               seasonal=FALSE,
                               stepwise=FALSE,
                               ic = c("aicc", "aic", "bic") ,
                               approximation=FALSE,
                               trace=FALSE)

# Forecast 12 observations
beerdat_arima_411_forecasts = forecast::forecast(beerdat_arima_411,h=12)
```

We can also produce the values into a nice table containing the forecasting points with confidence intervals

```
# Print the forecasts along with the boundaries
forecast_table<-print(beerdat_arima_411_forecasts) %>%
  mutate(observed=beer_dat_test,
         errors=`Point Forecast`-observed)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1990	5.000004	4.854874	5.145133	4.778047	5.221960
Apr 1990	5.004083	4.839863	5.168302	4.752931	5.255235
May 1990	5.028499	4.855981	5.201016	4.764656	5.292342
Jun 1990	5.087000	4.908593	5.265408	4.814150	5.359851
Jul 1990	5.121907	4.943000	5.300813	4.848293	5.395521
Aug 1990	5.142416	4.962550	5.322282	4.867334	5.417498
Sep 1990	5.149614	4.969024	5.330204	4.873425	5.425803
Oct 1990	5.136753	4.955477	5.318029	4.859515	5.413991
Nov 1990	5.121327	4.939987	5.302667	4.843992	5.398662
Dec 1990	5.107966	4.925893	5.290039	4.829510	5.386423
Jan 1991	5.099041	4.915647	5.282435	4.818564	5.379518
Feb 1991	5.099152	4.913904	5.284399	4.815840	5.382464

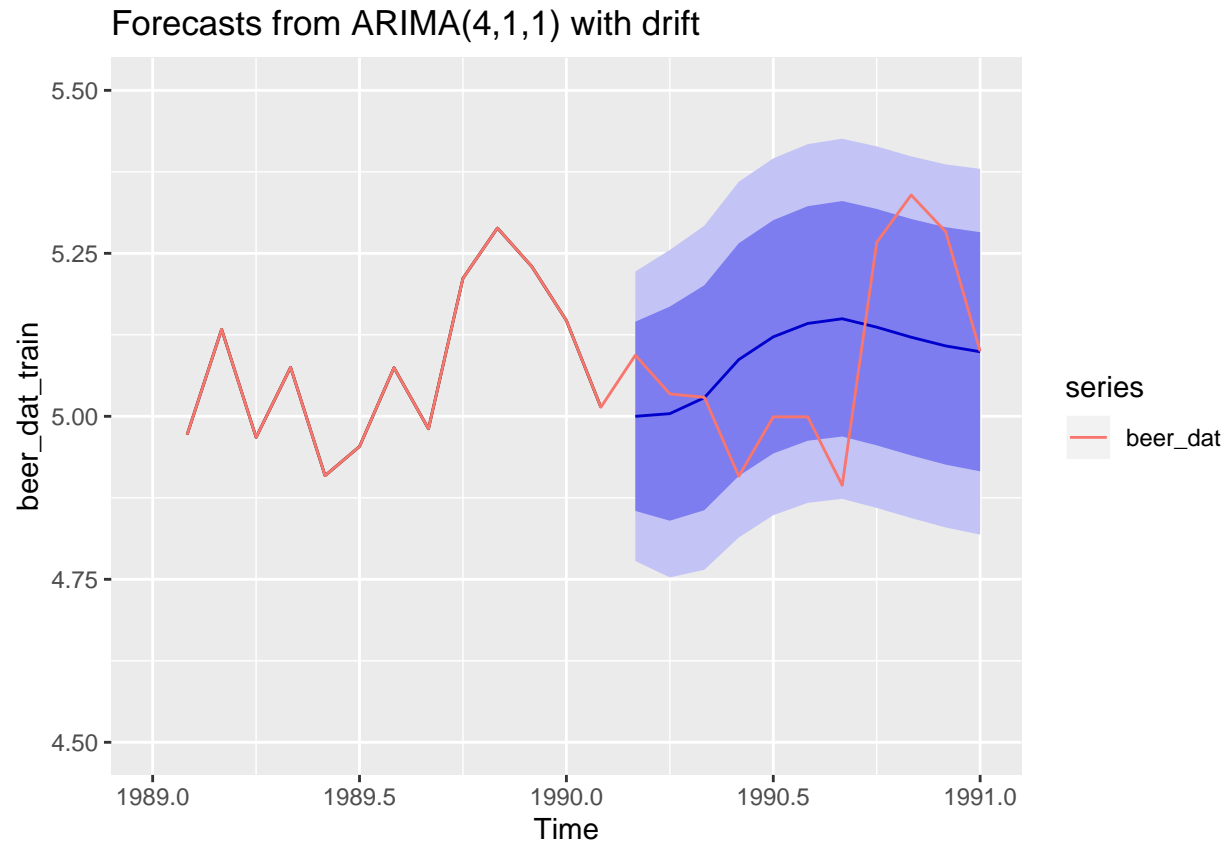
We see that although the model diagnostics were rather bad, all of the forecasted points still fall inside their respective confidence intervals

```
# Plot the predictions
autoplot(beerdat_arima_411_forecasts) +
  autolayer(beer_dat) + xlim(c(1989,1991)) + ylim(c(4.5,5.5))
```

Scale for 'x' is already present. Adding another scale for 'x', which will replace the existing scale.

Warning: Removed 397 rows containing missing values (geom_path).

Warning: Removed 398 rows containing missing values (geom_path).



Above: - The red is the original simulated data - Then the other is the best linear predictor for the data (blue line), along with confidence bounds (80% and 95%) - Most of the points are within the 80% confidence intervals

Evaluating predictions

We will now compare how many of these were “right”, by comparing them to the test data:

```
# Check how many points are inside the 95% confidence intervals
forecast_table %>%
  mutate(outside95=I(observed<`Lo 95` | observed>`Hi 95`)) %>%
  count(outside95)
```

```
# A tibble: 1 x 2
  outside95     n
  <I<lg1>> <int>
1 FALSE      12
```

```
forecast_table %>%
  mutate(outside80=I(observed<`Lo 80` | observed>`Hi 80`)) %>%
  count(outside80)
```

```
# A tibble: 2 x 2
  outside80     n
```

```
<I<lg1>> <int>
1 FALSE      9
2 TRUE       3
```

We observe that there are 3 points that fall outside the 80% confidence bounds, , whereas all of them fall within the 95% ones. This implies around 75% prediction accuracy.

Beer Data: Classical Decomposition approach

Background: Classical Decomposition

We saw that we had some problems with correctly estimating trend and seasonality for the **beer** data. We will now try the **classical decomposition** approach before fitting the series, i.e. we assume that for a time series $\{X_t\}$

$$X_t = \underbrace{m_t}_{trend} + \underbrace{s_t}_{season} + \underbrace{Y_t}_{noise}$$

where

$$\mathbb{E}[Y_t] = 0 \quad s_{t+d} = s_t, \forall t \quad \sum_{j=1}^d s_j = 0$$

So we will follow the following approach:

- 1) Estimate the trend \hat{m}_t , and remove it, i.e $\tilde{X}_t = X_t - \hat{m}_t \approx s_t + Y_t$ (by linear regression, moving avg, etc.)
- 2) Estimate the seasonal component, say \hat{s}_t , and obtain deseasonalized data $d_t = X_t - \hat{s}_t$, where

$$\hat{s}_k = w_k - \frac{1}{d} \sum_{i=1}^d w_i \quad w_k = \frac{1}{(t/d)} \sum_{j=0}^{t/d-1} (X_{k+jd} - \hat{m}_{k+jd})$$

- 3) Re-estimate the trend from d_t , say \tilde{m}_t , and calculate

$$\tilde{Y}_t = X_t - \hat{s}_t - \tilde{m}_t$$

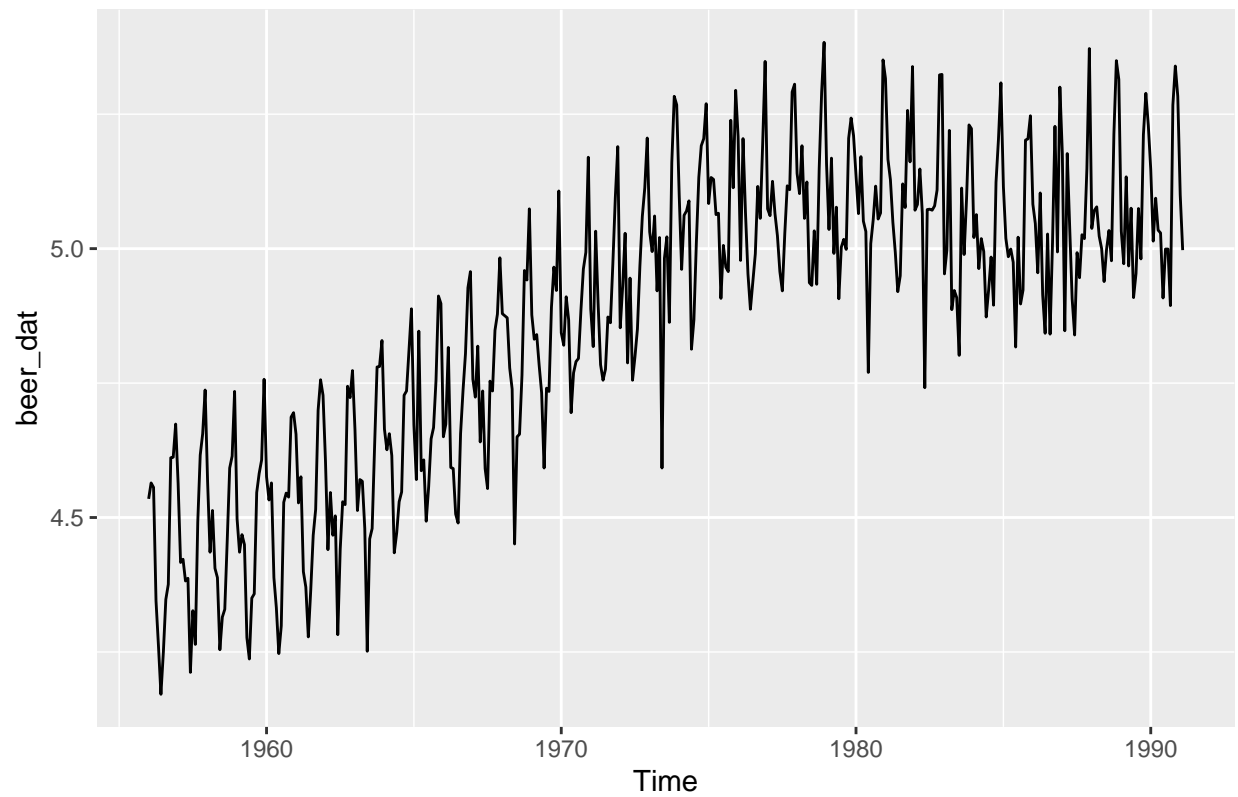
- 4) Finally, we will fit an ARMA model and make predictions on that! (then re-add the season and trend, and voilà)

Re-inspecting the data

Let's take a look at our data once more:

```
# Plot series
autoplot(beer_dat) + ggtitle("Beer data")
```

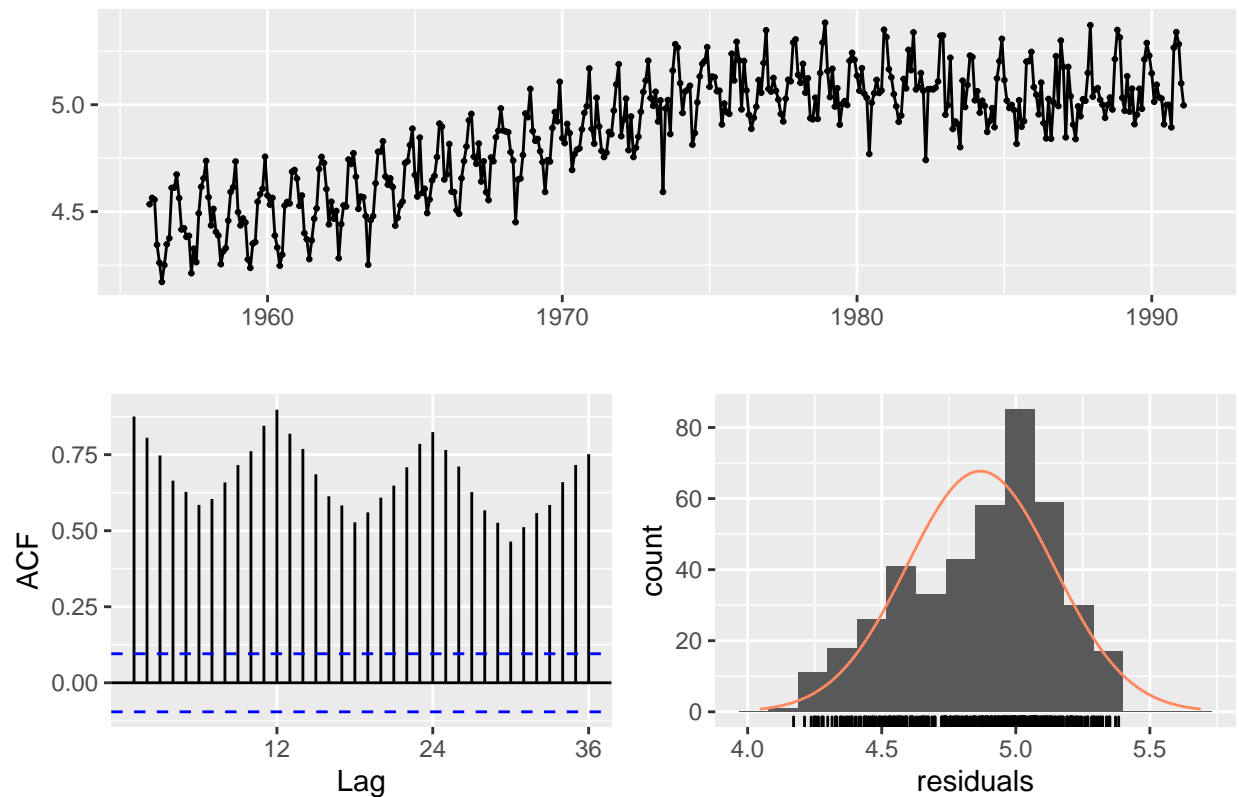
Beer data



```
checkresiduals(beer_dat)
```

Warning in modeldf.default(object): Could not find appropriate degrees of freedom for this model.

Residuals



We can observe once again there is a clear trend in the data, as well as a seasonal component present. The data is definitely not white noise.

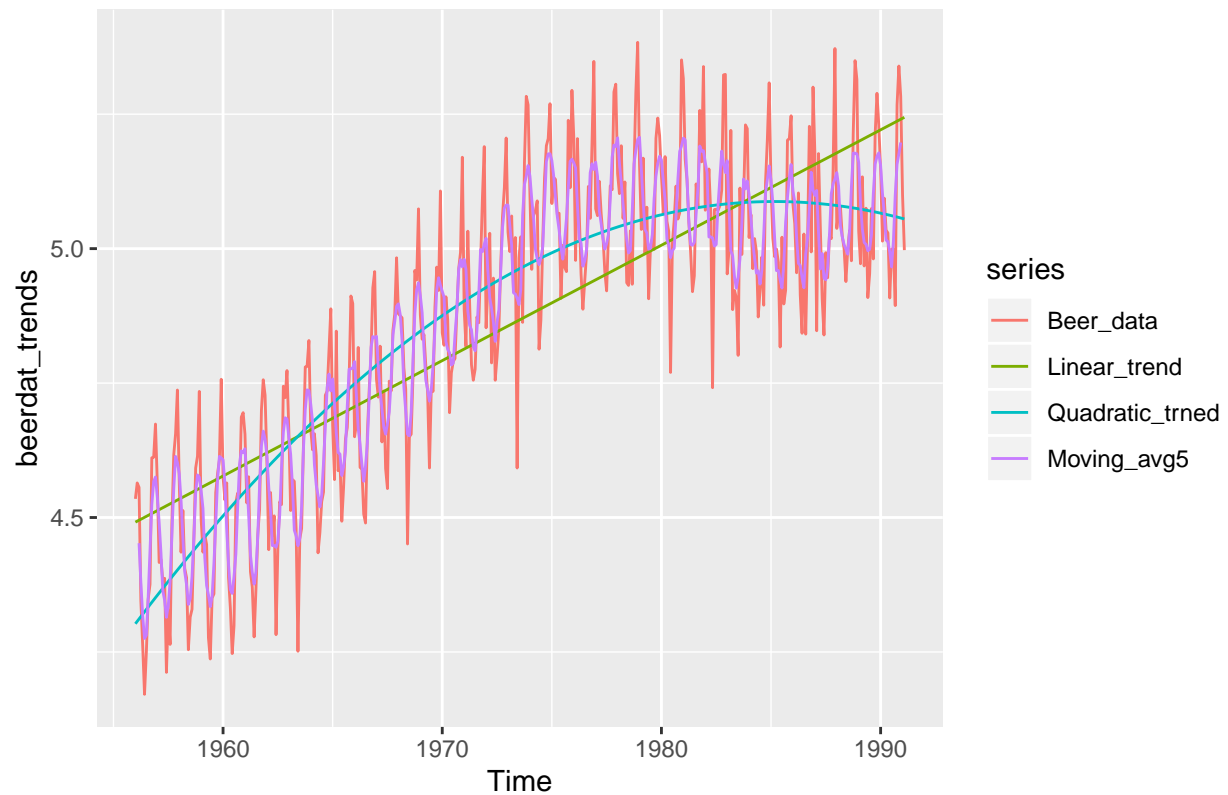
Estimating and removing trend

We will try estimating various trends to the data, being aware that if we over-estimate the trend we might lose predictive power later:

```
# Estimate various trends
beerdat_trend_linear <- tslm(beer_dat~trend)
beerdat_trend_quad <- tslm(beer_dat~trend + I(trend^2))
beerdat_trend_ma5 <- ma(beer_dat, order=5)
beerdat_trends <- cbind(Beer_data=beer_dat,
                        Linear_trend=fitted(beerdat_trend_linear),
                        Quadratic_trend=fitted(beerdat_trend_quad),
                        Moving_avg5 = beerdat_trend_ma5
                        )

# Plot all the trends together
autoplot(beerdat_trends) + ggtitle("Beer data trend estimations")
```

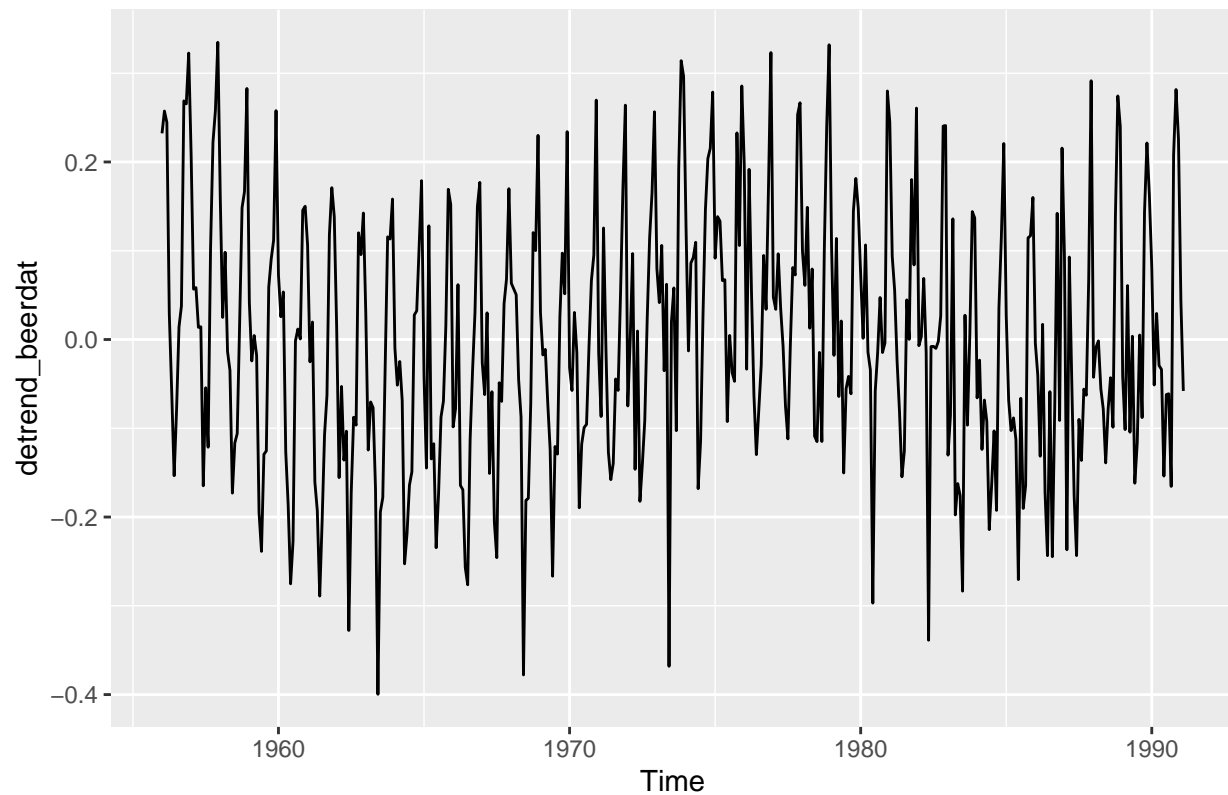
Beer data trend estimations



We will opt to remove the quadratic trend as the moving average may be a bit of an overkill :)

```
# Detrend and show the de-trended series
detrend_beerdat <- beer_dat - fitted(beerdat_trend_quad)
autoplot(detrend_beerdat) + ggtitle("De-trended Beer Data (quadratic trend)")
```

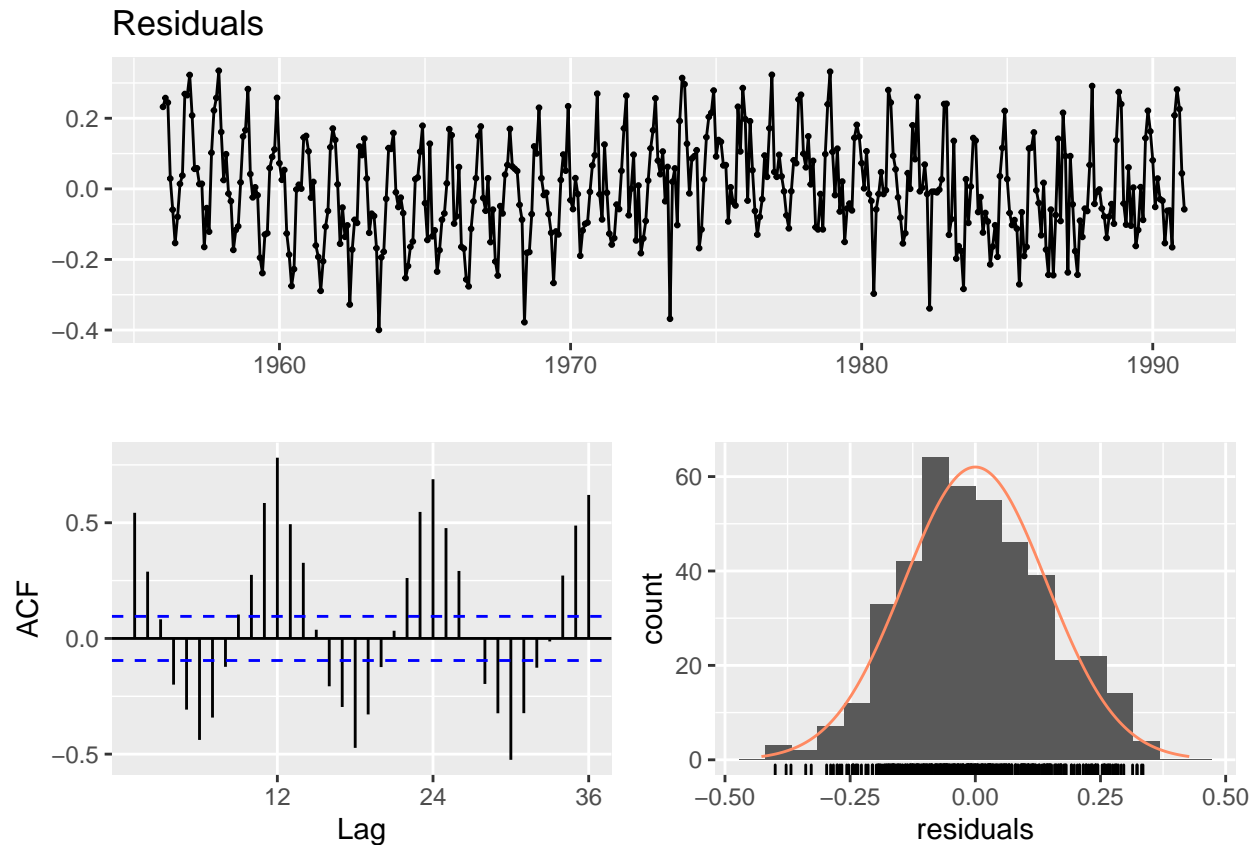

De-trended Beer Data (quadratic trend)



Let's quickly inspect the residuals

```
checkresiduals(detrend_beerdat)
```

```
Warning in modeldf.default(object): Could not find appropriate degrees of  
freedom for this model.
```

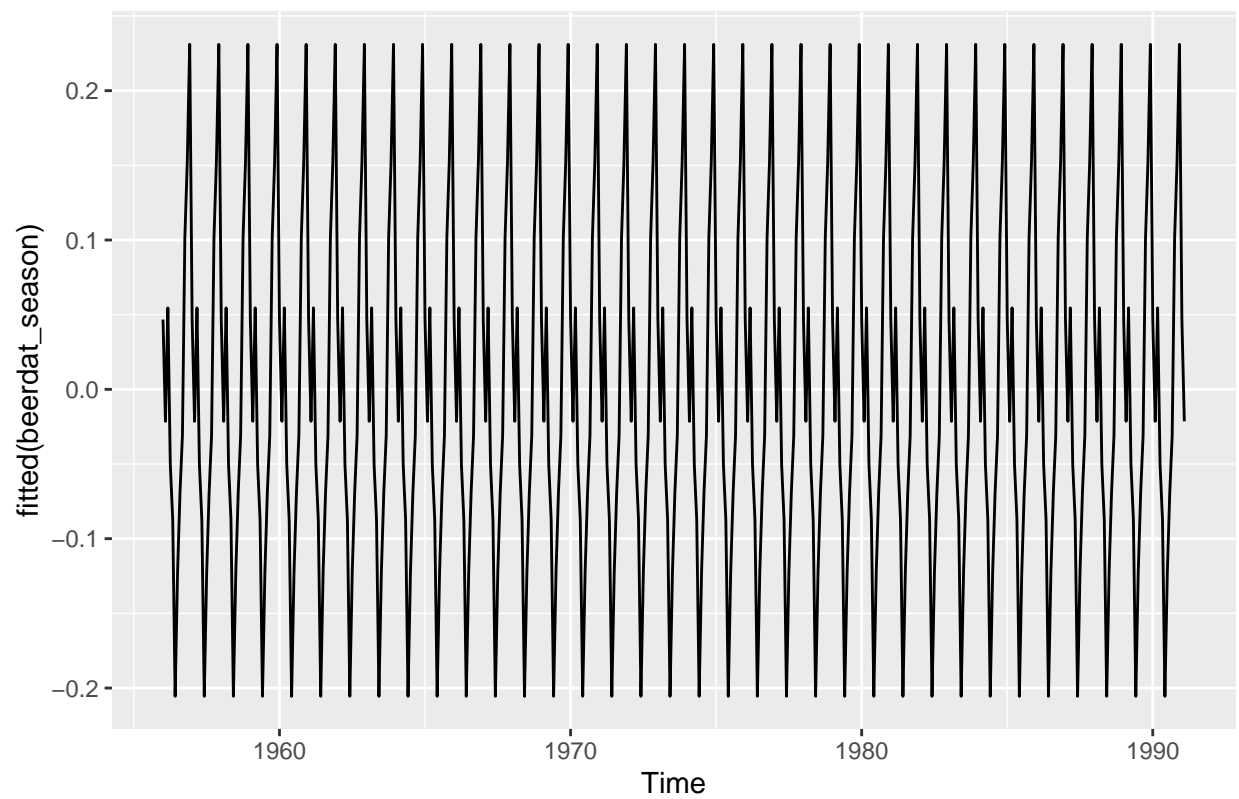


The residuals definitely seem a lot more centered towards 0, however there still seems to be a trend present (NOTE: it looks a bit like a sin/cos wave?) Let's however proceed with removing the seasonal component.

Removing seasonality

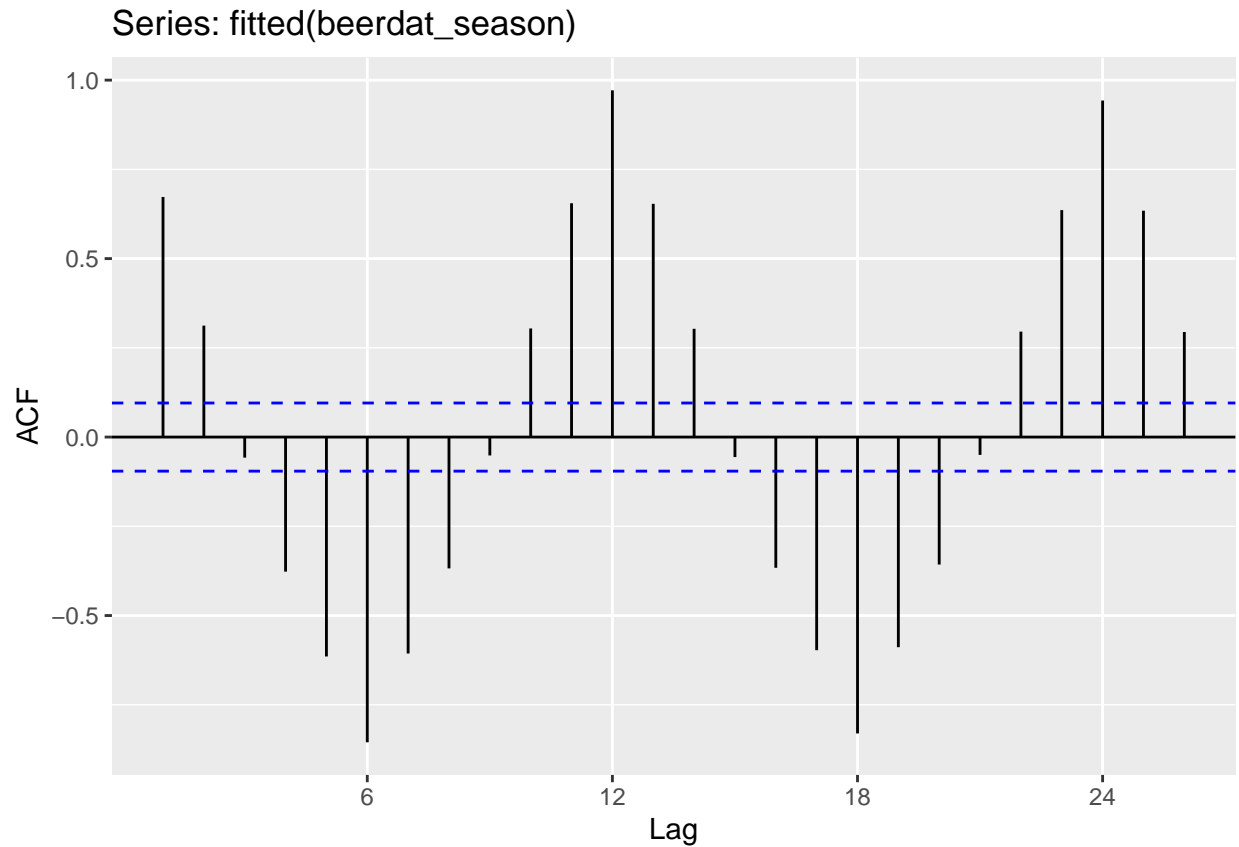
We will now estimate and remove the seasonal component.

```
# Estimate seasonality
beerdat_season <- tslm(detrend_beerdat ~ season)
autoplot(fitted(beerdat_season))
```



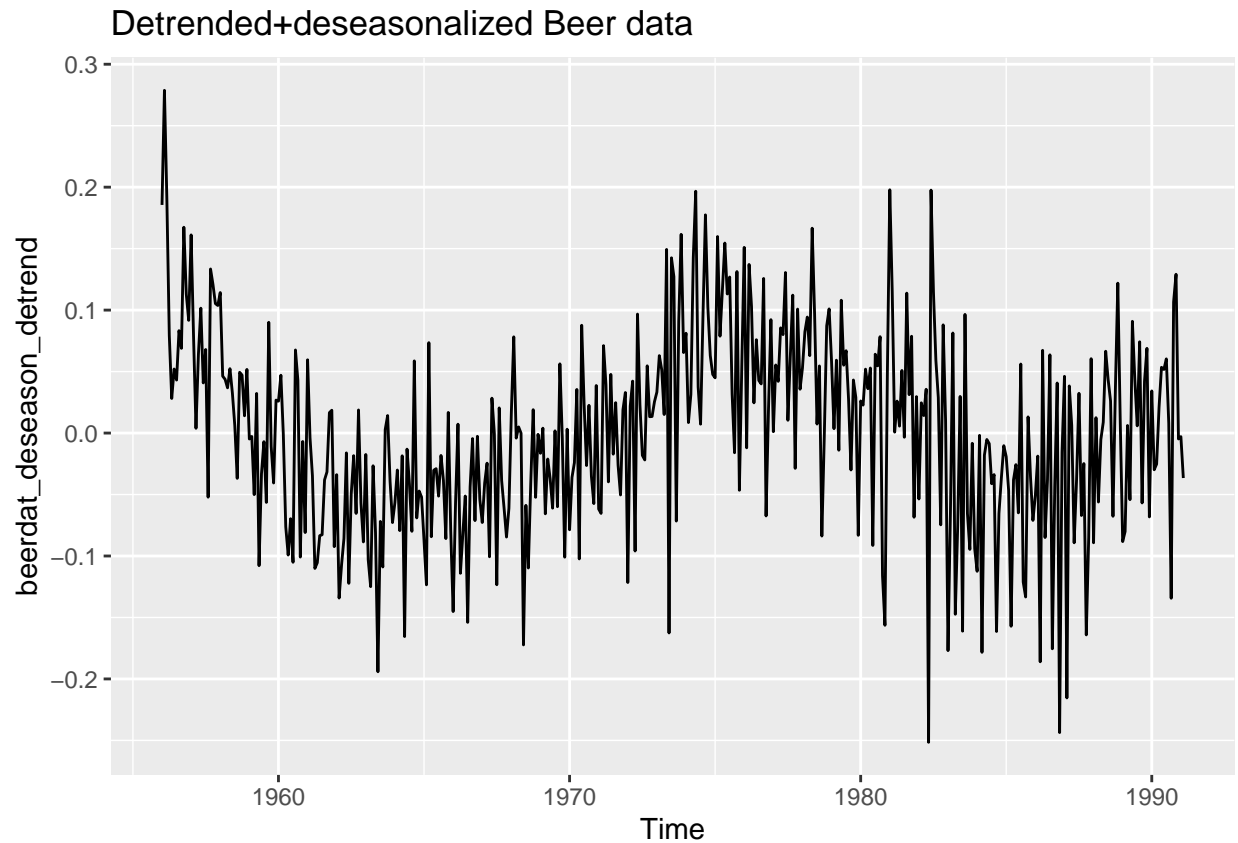
Great! Let's now take a look at the ACF

```
# ACF  
ggAcf(fitted(beerdat_season))
```



It seems that we have correctly estimated the seasonality, as the ACF clearly shows the seasonal dependency each 6 lags. Now, let's remove it from our data:

```
# Remove seasonal component and plot  
beerdat_deseason_detrend <- (detrend_beerdat - fitted(beerdat_season))  
autoplot(beerdat_deseason_detrend) + ggtitle("Detrended+deseasonalized Beer data")
```



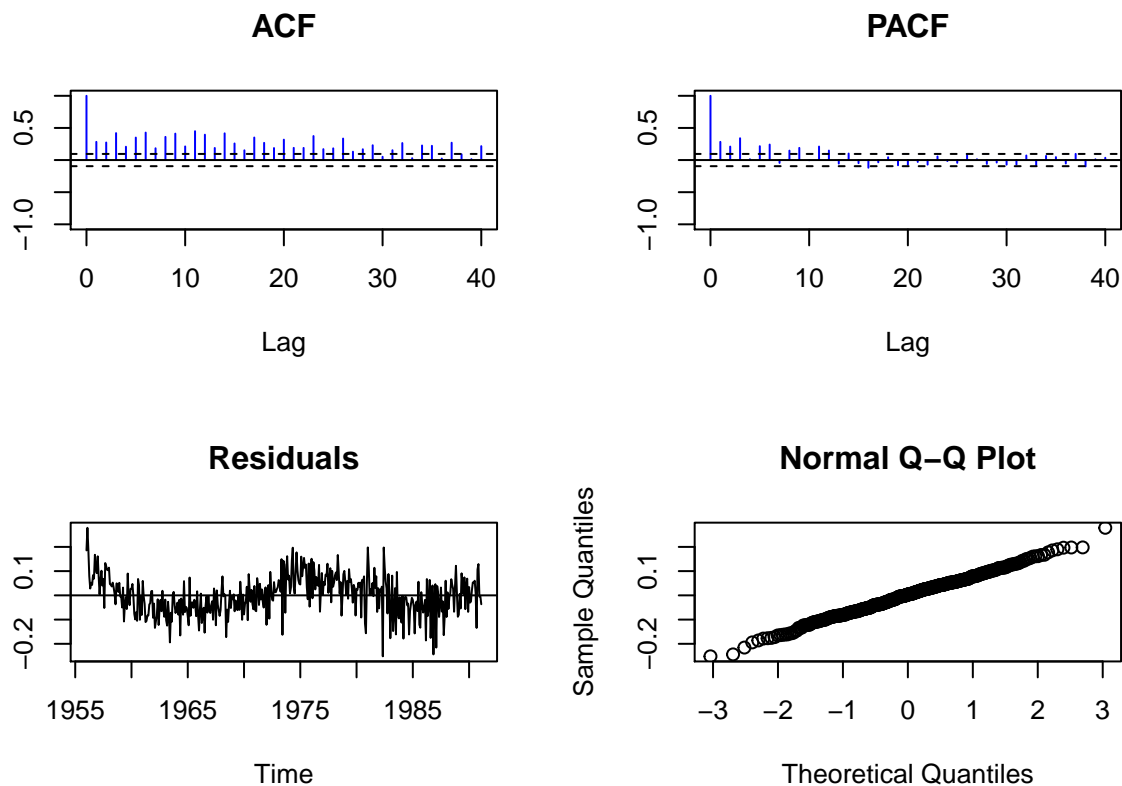
OH boy, that doesn't look like white noise...

Testing the detrended + deseasonalized data

```
test(beerdats_deseason_detrend)
```

Null hypothesis: Residuals are iid noise.

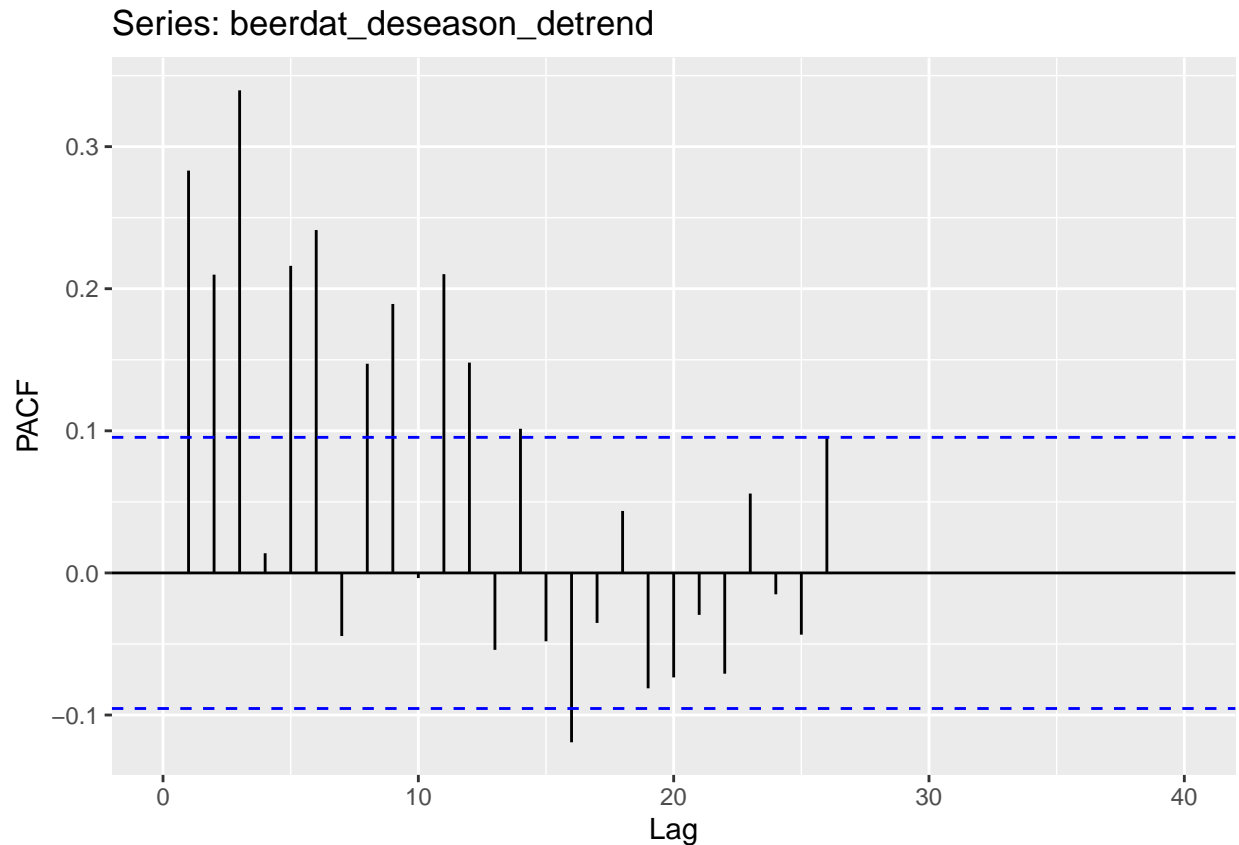
Test	Distribution	Statistic	p-value
Ljung-Box Q	$Q \sim \text{chisq}(20)$	887.9	0 *
McLeod-Li Q	$Q \sim \text{chisq}(20)$	72.55	0 *
Turning points T	$(T-280)/8.6 \sim N(0,1)$	287	0.418
Diff signs S	$(S-210.5)/5.9 \sim N(0,1)$	205	0.3543
Rank P	$(P-44415.5)/1447.4 \sim N(0,1)$	45222	0.5774



Funnily enough, although the residuals definitely don't look like white noise, and the series seems not to be stationary according to three out of the five tests (however note that the Ljung-Box and the McLeod-Li seem to indicate the opposite), we can see from the PACF that in fact most of the correlation can be explained by the first ~12 lags, which is indicative of the seasonal component:

```
# Plot PACF for residuals
ggPacf(beerdat_deseason_detrend) + xlim(c(0,40))
```

Scale for 'x' is already present. Adding another scale for 'x', which will replace the existing scale.



Fitting an ARMA model

We will still try to fit an ARMA model to these series, and see if they perform better than the ARIMA estimation we performed before. However, we will have to separate the data into train and test sets since we detrended and deseasonalized the whole original data.

```
# Preprocess data
beer_dat_train_dd <- window(beerdat_deseason_detrend, end=1990.1) # train dataset
beer_dat_test_dd <- window(beerdat_deseason_detrend, start=1990.1) # test dataset
str(beer_dat_train_dd)
```

```
Time-Series [1:410] from 1956 to 1990: 0.1855 0.2788 0.19 0.0795 0.0281 ...
```

```
str(beer_dat_test_dd)
```

```
Time-Series [1:12] from 1990 to 1991: -0.0252 0.0209 0.0535 0.0517 0.0603 ...
```

```
# Fit the ARIMA model on train data
beerdat_arma = auto.arima(beer_dat_train_dd ,
  seasonal=FALSE,
  stepwise=FALSE,
  max.d= 0, # don't allow ARIMA
  ic = c("aicc", "aic", "bic") ,
  approximation=FALSE,
  trace=TRUE)
```

ARIMA(0,0,0)	with zero mean	: -896.8996
ARIMA(0,0,0)	with non-zero mean	: -894.8991
ARIMA(0,0,1)	with zero mean	: -921.0555
ARIMA(0,0,1)	with non-zero mean	: -919.037
ARIMA(0,0,2)	with zero mean	: -928.4903
ARIMA(0,0,2)	with non-zero mean	: -926.455
ARIMA(0,0,3)	with zero mean	: -965.251
ARIMA(0,0,3)	with non-zero mean	: -963.2019
ARIMA(0,0,4)	with zero mean	: -964.4386
ARIMA(0,0,4)	with non-zero mean	: -962.3789
ARIMA(0,0,5)	with zero mean	: -979.6284
ARIMA(0,0,5)	with non-zero mean	: -977.5614
ARIMA(1,0,0)	with zero mean	: -930.5653
ARIMA(1,0,0)	with non-zero mean	: -928.5411
ARIMA(1,0,1)	with zero mean	: Inf
ARIMA(1,0,1)	with non-zero mean	: Inf
ARIMA(1,0,2)	with zero mean	: Inf
ARIMA(1,0,2)	with non-zero mean	: Inf
ARIMA(1,0,3)	with zero mean	: Inf
ARIMA(1,0,3)	with non-zero mean	: Inf
ARIMA(1,0,4)	with zero mean	: Inf
ARIMA(1,0,4)	with non-zero mean	: Inf
ARIMA(2,0,0)	with zero mean	: -951.3139
ARIMA(2,0,0)	with non-zero mean	: -949.2745
ARIMA(2,0,1)	with zero mean	: Inf
ARIMA(2,0,1)	with non-zero mean	: Inf
ARIMA(2,0,2)	with zero mean	: Inf
ARIMA(2,0,2)	with non-zero mean	: Inf
ARIMA(2,0,3)	with zero mean	: Inf
ARIMA(2,0,3)	with non-zero mean	: Inf
ARIMA(3,0,0)	with zero mean	: -1006.679
ARIMA(3,0,0)	with non-zero mean	: -1004.657
ARIMA(3,0,1)	with zero mean	: Inf
ARIMA(3,0,1)	with non-zero mean	: Inf
ARIMA(3,0,2)	with zero mean	: Inf
ARIMA(3,0,2)	with non-zero mean	: Inf
ARIMA(4,0,0)	with zero mean	: -1004.848
ARIMA(4,0,0)	with non-zero mean	: -1002.82
ARIMA(4,0,1)	with zero mean	: Inf
ARIMA(4,0,1)	with non-zero mean	: Inf
ARIMA(5,0,0)	with zero mean	: -1021.291
ARIMA(5,0,0)	with non-zero mean	: -1019.29

Best model: ARIMA(5,0,0) with zero mean

So according to the `auto.arima`, our best model is an AR(5) model. Let's inspect it.

Inspecting the chosen AR(5) model

Now let's inspect our model:


```
beerdat_ar5 <- beerdat_arma
beerdat_ar5
```

```
Series: beer_dat_train_dd
ARIMA(5,0,0) with zero mean
```

```
Coefficients:
```

```
      ar1      ar2      ar3      ar4      ar5
      0.1284 0.0809 0.3327 0.0008 0.2143
s.e. 0.0483 0.0493 0.0466 0.0494 0.0492
```

```
sigma^2 estimated as 0.004755: log likelihood=516.75
AIC=-1021.5 AICc=-1021.29 BIC=-997.4
```

This says that we have a lag 1 difference giving an ARMA(3,1) model. From this, we see all the estimated parameters $\hat{\phi}_1, \hat{\phi}_2, \hat{\phi}_3, \hat{\phi}_4, \hat{\phi}_5$ and $\hat{\sigma}^2$. However we can see that the coefficient $\hat{\phi}_4$ is pretty much zero, so we could assume this is in fact an AR(4) model.

Parameters

We can construct 95% confidence intervals for these:

```
## Calculate the wald confidence interval
params_ar5 <- beerdat_ar5$coef # extract parameters
se_ar5 <- sqrt(diag(beerdat_ar5$var.coef)) # extract s.e. for each
z <- qnorm(0.975) # this is the alpha/2 quantile
c.upper <- params_ar5 + z*se_ar5 # upper bound
c.lower <- params_ar5 - z*se_ar5 # lower bound
CI_ar5 <- cbind(params_ar5, se_ar5, c.lower, c.upper)
colnames(CI_ar5) <- c("parameter", "s.d.", "2.5 %", "97.5 %")
CI_ar5
```

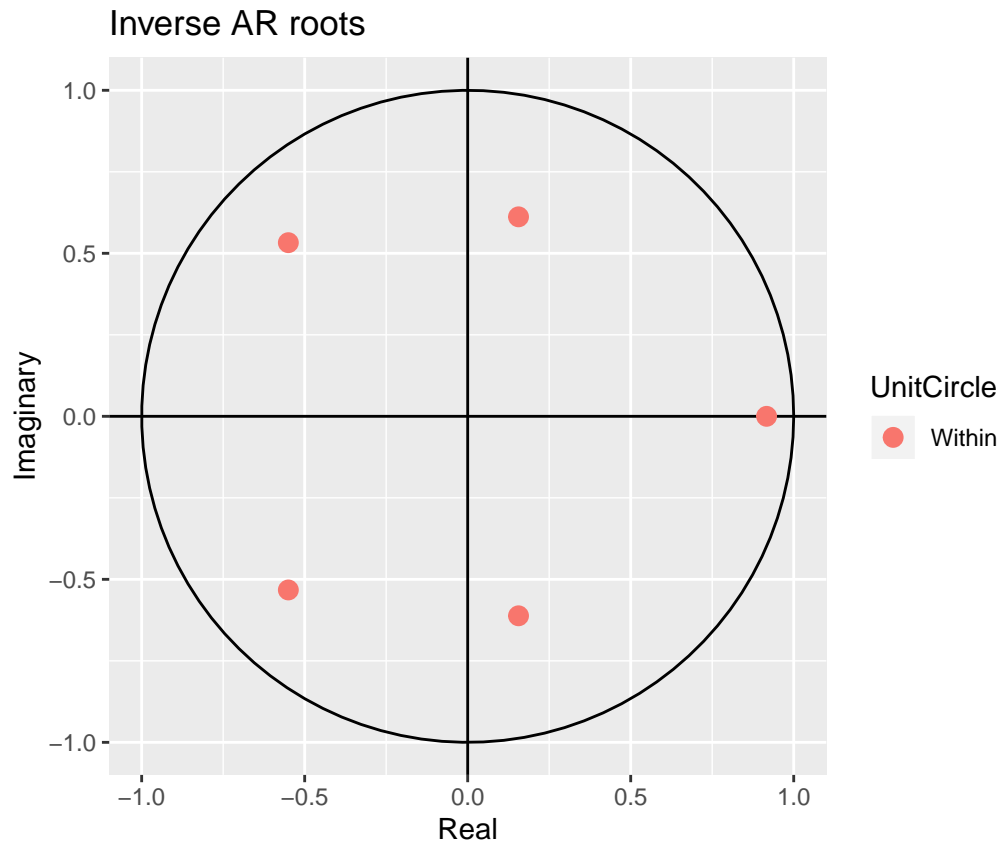
```
      parameter      s.d.      2.5 %      97.5 %
ar1 0.1284392295 0.04829561 0.03378157 0.2230969
ar2 0.0808822273 0.04925638 -0.01565850 0.1774229
ar3 0.3326715559 0.04659216 0.24135259 0.4239905
ar4 0.0007815309 0.04938875 -0.09601864 0.0975817
ar5 0.2142980228 0.04920953 0.11784912 0.3107469
```

Again, note that the ar4 coefficient is estimated to be zero.

Residuals inspection and diagnostics

Let's first verify the roots of the selected AR(5) process:

```
# Plot the toods
autoplot(beerdat_ar5)
```

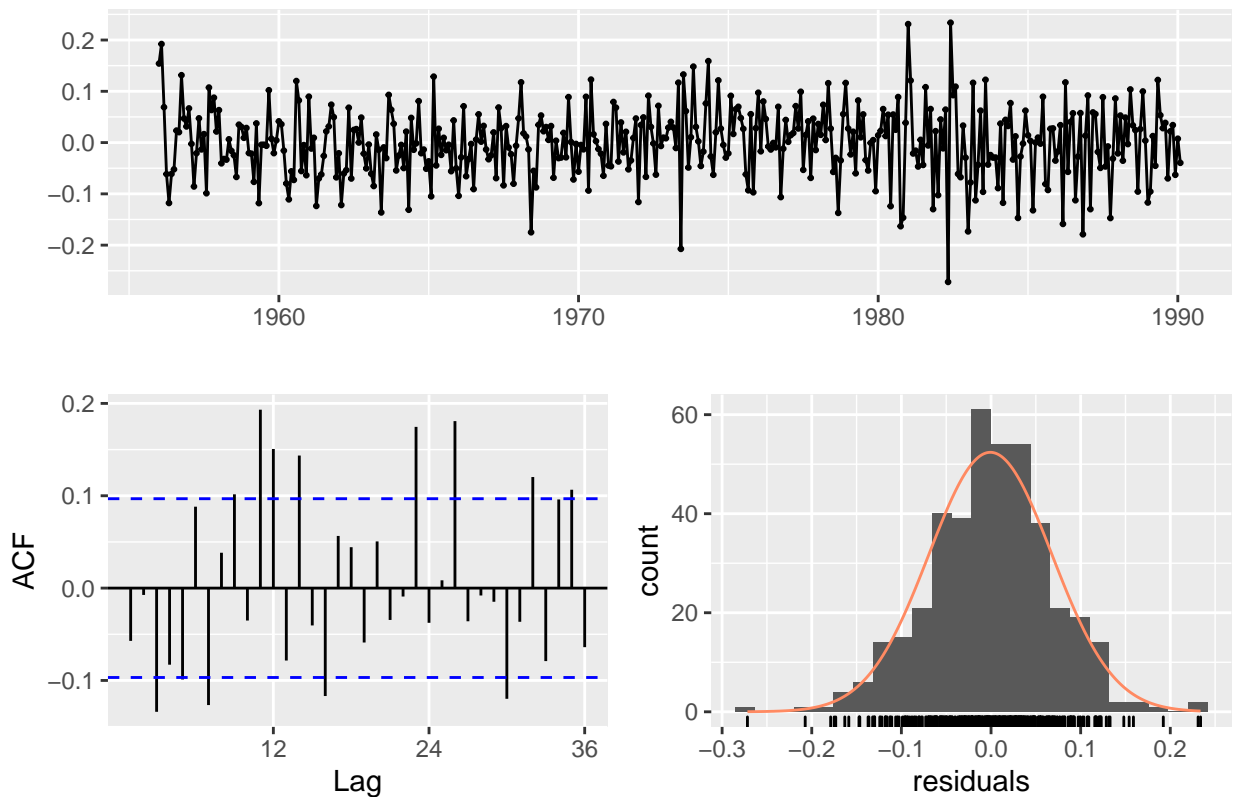


We see that they are all between, so for sure the process is both stationary and causal. The root that is very close to boundary seems to be the one for the ar4 coefficient.

Let's observe the residuals from the resulting model:

```
checkresiduals(beerdat_ar5)
```

Residuals from ARIMA(5,0,0) with zero mean



Ljung-Box test

```
data: Residuals from ARIMA(5,0,0) with zero mean
Q* = 93.909, df = 19, p-value = 6.682e-12
```

```
Model df: 5. Total lags used: 24
```

According to the Ljung-Box test, we reject the hypothesis of stationarity, as we would expect. However note that the roots tells us otherwise!

```
# Test with a bunch of different k's ? (bigger augmented versions)
adf.test(residuals(beerdat_ar5))
```

```
Warning in adf.test(residuals(beerdat_ar5)): p-value smaller than printed p-
value
```

Augmented Dickey-Fuller Test

```
data: residuals(beerdat_ar5)
Dickey-Fuller = -9.1761, Lag order = 7, p-value = 0.01
alternative hypothesis: stationary
```

```
kpss.test(residuals(beerdat_ar5))
```

Warning in `kpss.test(residuals(beerdat_ar5))`: p-value greater than printed p-value

KPSS Test for Level Stationarity

data: residuals(beerdat_ar5)

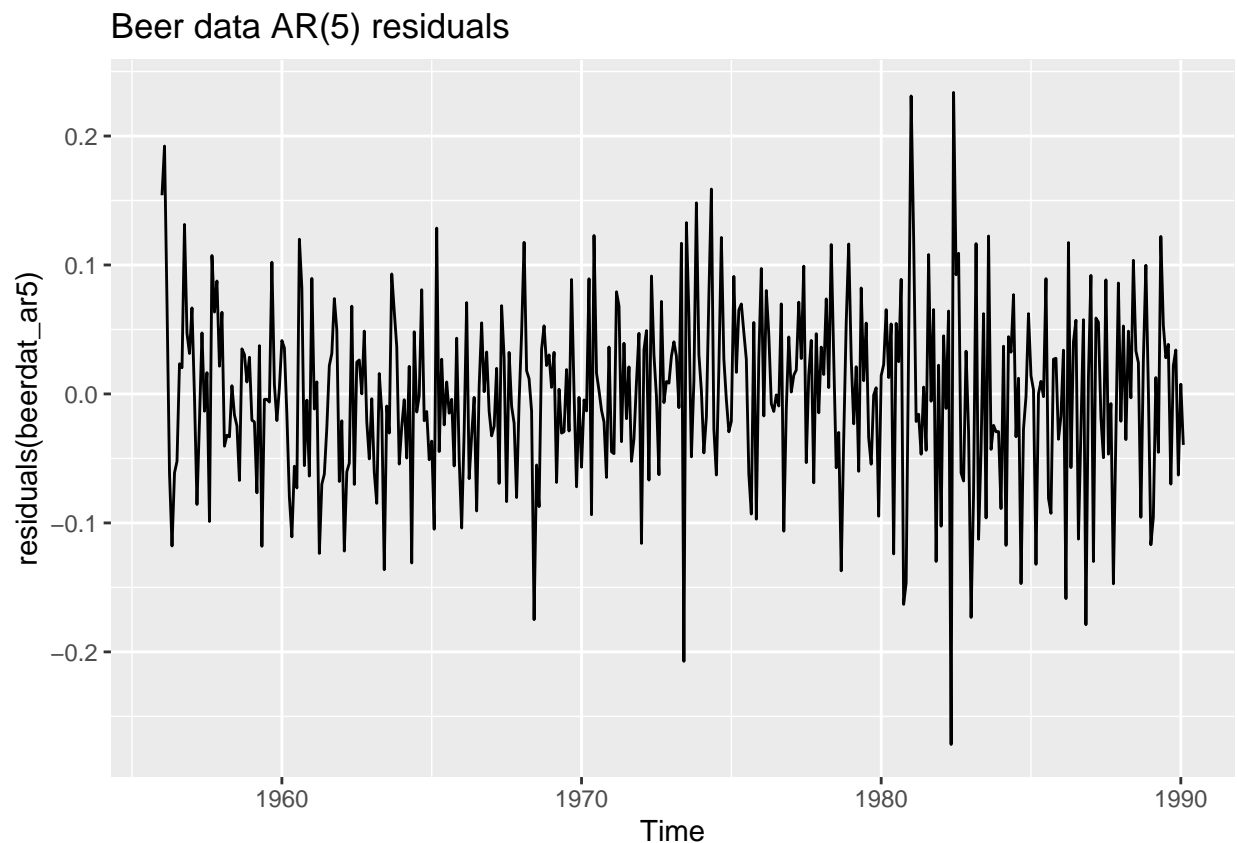
KPSS Level = 0.19311, Truncation lag parameter = 5, p-value = 0.1

- The adf test indicates a stationary process
- The kpss test also indicates a stationary process, as we fail to reject the null.

Eureka! It seems that this process is indeed stationary.

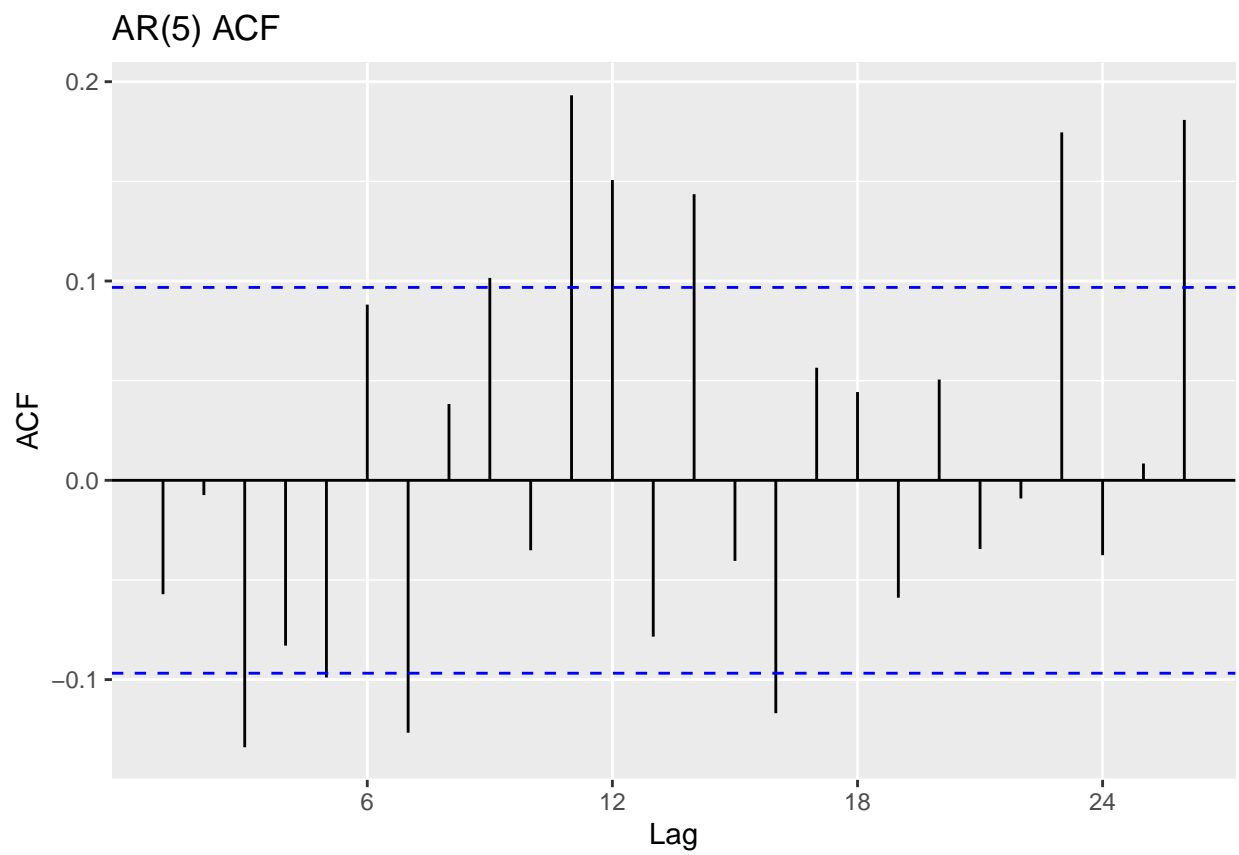
```
# Plot the resulting model
```

```
autoplot(residuals(beerdat_ar5)) + ggtitle("Beer data AR(5) residuals")
```

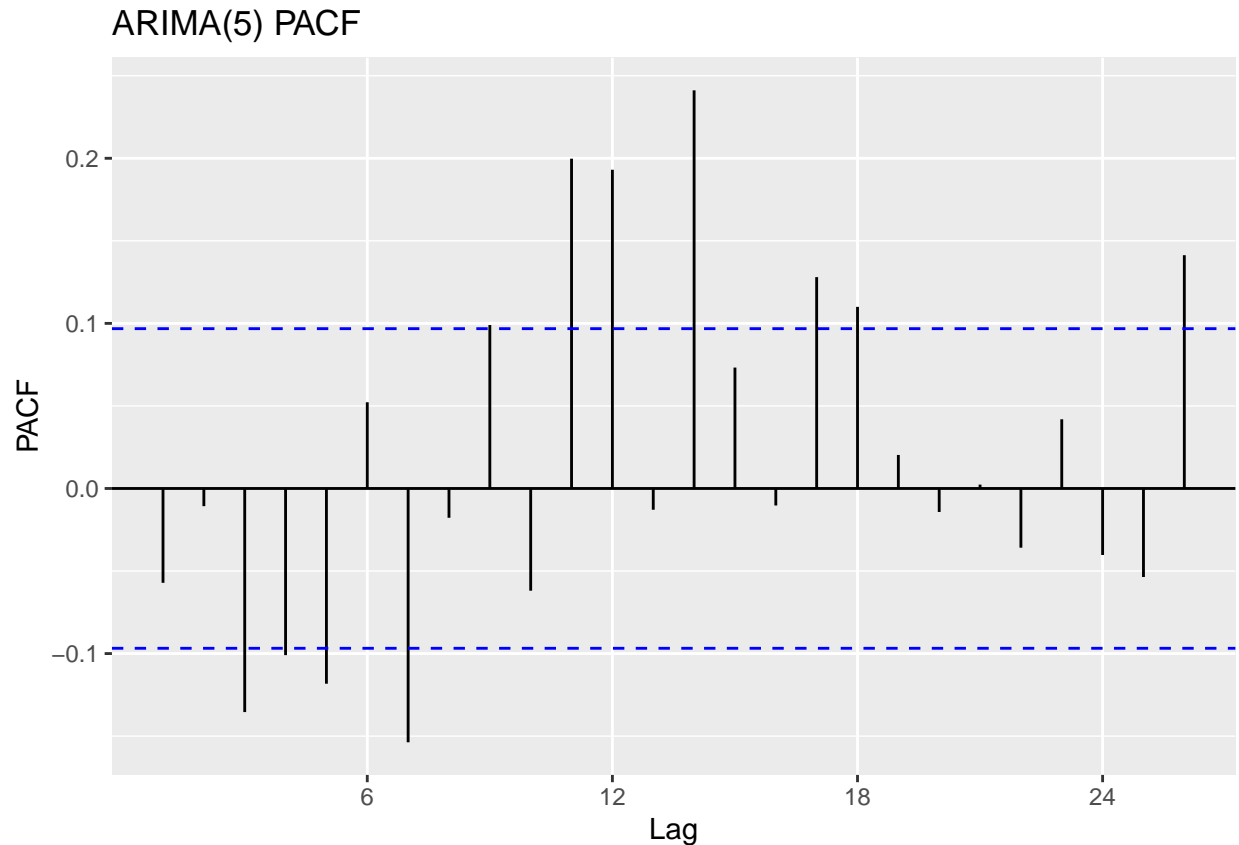


We observe that indeed the model seems to be a lot more centered towards 0, and seems a little bit more like white noise. Let us check this with the ACF and PACF plots:

```
# Plot ACF
ggAcf(residuals(beerdat_ar5)) + ggtitle("AR(5) ACF")
```



```
# Plot PACF
ggPacf(residuals(beerdat_ar5)) + ggtitle("ARIMA(5) PACF")
```



Again, some of the residuals go out bounds, but the guess is that we **could** have estimated better the average (for example with a moving average instead of a quadratic trend).

Forecasting

Just for clarity, let's re-fit the ARIMA(5) model with the training data and predict on the test data. However, since we fitted the AR(5) with detrended and deseasonalized data, we will also detrend and deseasonalize the test data

```
# Fit the ARMA model on train data
beerdat_ar5 = auto.arima(beer_dat_train_dd,
  seasonal=FALSE,
  stepwise=FALSE,
  max.d= 0, # don't allow ARIMA
  ic = c("aicc", "aic", "bic") ,
  approximation=FALSE,
  trace=FALSE)

# Forecast ar5 predictions, trend and seasonality
beerdat_ar5_forecasts <- forecast::forecast(beerdat_ar5,h=12) # AR(5) forecasts
forecasted_trend <- forecast::forecast(beerdat_trend_quad, h=12) # trend
forecasted_season <- forecast::forecast(beerdat_season, h=12) # season
```

Let's see a plot of the deseasoned+detrended data and predictions:

```
# Print the forecasts along with the boundaries
forecast_table<-print(beerdat_ar5_forecasts) %>%
  mutate(observed=beer_dat_test_dd,
         errors=`Point Forecast`-observed)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1990	-0.0146694895	-0.10303931	0.07370033	-0.1498195	0.1204805
Apr 1990	0.0218086675	-0.06728707	0.11090441	-0.1144515	0.1580688
May 1990	-0.0229049410	-0.11241529	0.06660541	-0.1597992	0.1139893
Jun 1990	0.0012585370	-0.09360668	0.09612375	-0.1438253	0.1463424
Jul 1990	-0.0008442878	-0.09609642	0.09440784	-0.1465199	0.1448313
Aug 1990	-0.0107530672	-0.10934708	0.08784094	-0.1615396	0.1400335
Sep 1990	0.0036249293	-0.09639938	0.10364924	-0.1493491	0.1565989
Oct 1990	-0.0055925194	-0.10601603	0.09483100	-0.1591770	0.1479920
Nov 1990	-0.0037333040	-0.10570814	0.09824154	-0.1596904	0.1522238
Dec 1990	0.0000847397	-0.10242561	0.10259509	-0.1566913	0.1568608
Jan 1991	-0.0044530742	-0.10757569	0.09866954	-0.1621655	0.1532594
Feb 1991	-0.0010346151	-0.10483030	0.10276107	-0.1597764	0.1577072

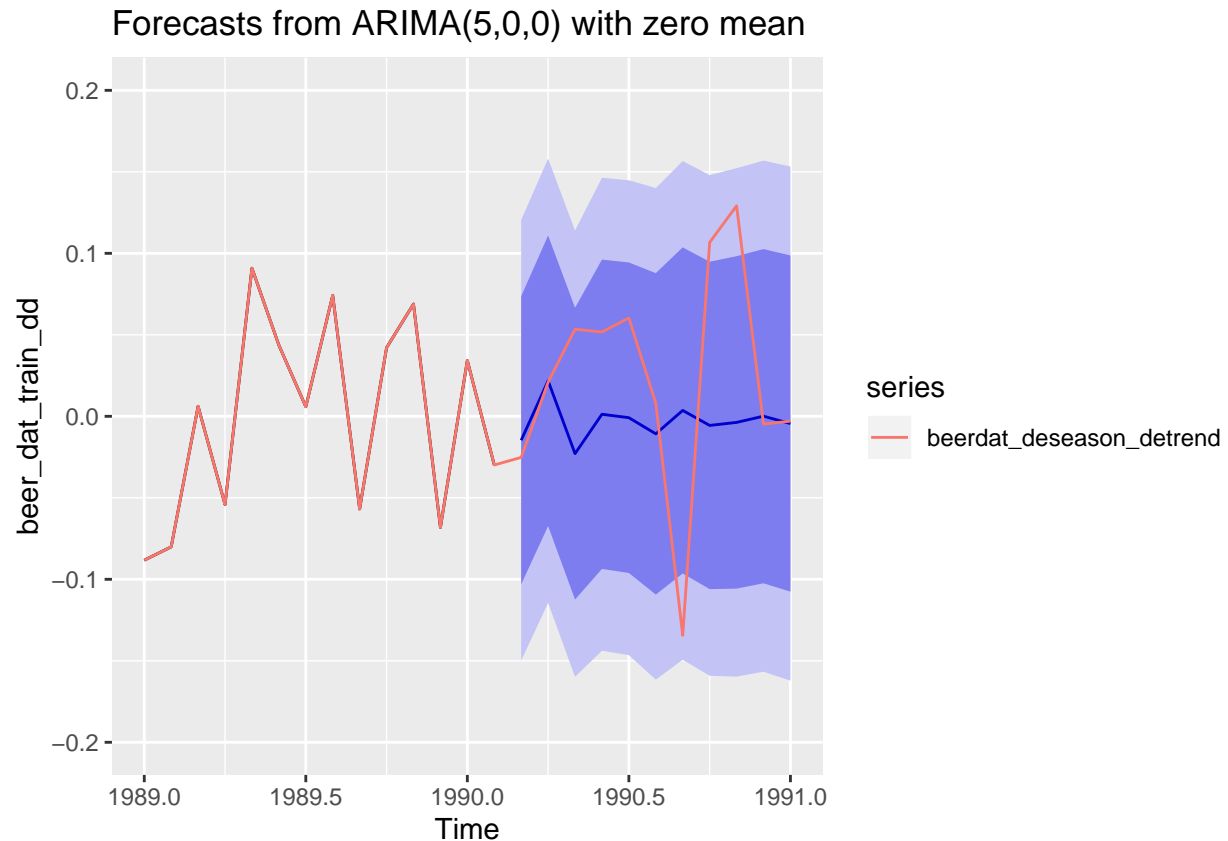
Now we can plot the forecasted points!

```
# Plot the predictions
autoplot(beerdat_ar5_forecasts) +
  autolayer(beerdat_deseason_detrend) + xlim(c(1989,1991)) + ylim(c(-0.2,0.2))
```

Scale for 'x' is already present. Adding another scale for 'x', which will replace the existing scale.

Warning: Removed 396 rows containing missing values (geom_path).

Warning: Removed 397 rows containing missing values (geom_path).



Above: - The red is the original simulated data - Then the other is the best linear predictor for the data (blue line), along with confidence bounds (80% and 95%) - Most of the points are within the 80% confidence intervals

Evaluating predictions

We will now compare how many of these were “right”, by comparing them to the test data:

```
# Check how many points are inside the 95% confidence intervals
forecast_table %>%
  mutate(outside95=I(observed<`Lo 95` | observed>`Hi 95`)) %>%
  count(outside95)
```

```
# A tibble: 1 x 2
  outside95     n
  <I<lg1>> <int>
1 FALSE      12
```

```
forecast_table %>%
  mutate(outside80=I(observed<`Lo 80` | observed>`Hi 80`)) %>%
  count(outside80)
```

```
# A tibble: 2 x 2
  outside80     n
```



```

<I<lg1>> <int>
1 FALSE      9
2 TRUE       3

```

Just as before, we obtain the same amount of points outside the interval, indicating that the predictive power of both models is roughly the same.

Aside: Restoring the original scale

- **NOTE:** If we were to convert these values back to the original scale, below would be the beginning of the procedure:

Save each forecast table to access individual values and alter the test data:

```

# Save data.frame objects
forecast_tbl <- print(beerdat_ar5_forecasts) # ar(5)

```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1990	-0.0146694895	-0.10303931	0.07370033	-0.1498195	0.1204805
Apr 1990	0.0218086675	-0.06728707	0.11090441	-0.1144515	0.1580688
May 1990	-0.0229049410	-0.11241529	0.06660541	-0.1597992	0.1139893
Jun 1990	0.0012585370	-0.09360668	0.09612375	-0.1438253	0.1463424
Jul 1990	-0.0008442878	-0.09609642	0.09440784	-0.1465199	0.1448313
Aug 1990	-0.0107530672	-0.10934708	0.08784094	-0.1615396	0.1400335
Sep 1990	0.0036249293	-0.09639938	0.10364924	-0.1493491	0.1565989
Oct 1990	-0.0055925194	-0.10601603	0.09483100	-0.1591770	0.1479920
Nov 1990	-0.0037333040	-0.10570814	0.09824154	-0.1596904	0.1522238
Dec 1990	0.0000847397	-0.10242561	0.10259509	-0.1566913	0.1568608
Jan 1991	-0.0044530742	-0.10757569	0.09866954	-0.1621655	0.1532594
Feb 1991	-0.0010346151	-0.10483030	0.10276107	-0.1597764	0.1577072

```

forecast_trend_tbl <- print(forecasted_trend) # trend

```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1991	5.054253	4.869076	5.239430	4.770676	5.337830
Apr 1991	5.053322	4.868108	5.238537	4.769688	5.336956
May 1991	5.052379	4.867126	5.237631	4.768687	5.336071
Jun 1991	5.051423	4.866132	5.236713	4.767672	5.335173
Jul 1991	5.050453	4.865123	5.235783	4.766643	5.334264
Aug 1991	5.049471	4.864102	5.234841	4.765600	5.333343
Sep 1991	5.048477	4.863067	5.233887	4.764543	5.332410
Oct 1991	5.047469	4.862018	5.232920	4.763473	5.331465
Nov 1991	5.046449	4.860956	5.231941	4.762389	5.330509
Dec 1991	5.045415	4.859880	5.230950	4.761291	5.329540
Jan 1992	5.044369	4.858791	5.229947	4.760179	5.328560
Feb 1992	5.043310	4.857689	5.228932	4.759053	5.327568

```

forecast_season_tbl <- print(forecasted_season) # season

```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1991	0.05454736	-0.051833321	0.16092804	-0.10836662	0.21746135

Apr 1991	-0.05022896	-0.156609638	0.05615173	-0.21314294	0.11268503
May 1991	-0.08718772	-0.193568399	0.01919297	-0.25010170	0.07572627
Jun 1991	-0.20553507	-0.311915749	-0.09915438	-0.36844905	-0.04262108
Jul 1991	-0.12233736	-0.228718040	-0.01595667	-0.28525134	0.04057663
Aug 1991	-0.06916555	-0.175546232	0.03721513	-0.23207954	0.09374844
Sep 1991	-0.03109403	-0.137474711	0.07528665	-0.19400801	0.13181996
Oct 1991	0.10142187	-0.004958816	0.20780255	-0.06149212	0.26433585
Nov 1991	0.15246770	0.046087018	0.25884838	-0.01044628	0.31538169
Dec 1991	0.23103190	0.124651221	0.33741259	0.06811792	0.39394589
Jan 1992	0.04672604	-0.059613598	0.15306567	-0.11612509	0.20957716
Feb 1992	-0.02137064	-0.127710269	0.08496900	-0.18422176	0.14148048

```
total_forecast <- forecast_tbl + forecast_trend_tbl + forecast_season_tbl
total_forecast
```

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Mar 1990		5.094131	4.714203	5.474059	4.512490	5.675772
Apr 1990		5.024902	4.644211	5.405593	4.442094	5.607710
May 1990		4.942286	4.561143	5.323430	4.358786	5.525786
Jun 1990		4.847146	4.460609	5.233683	4.255397	5.438895
Jul 1990		4.927272	4.540309	5.314235	4.334872	5.519672
Aug 1990		4.969553	4.579208	5.359897	4.371981	5.567125
Sep 1990		5.021008	4.629193	5.412823	4.421186	5.620829
Oct 1990		5.143298	4.751043	5.535554	4.542804	5.743793
Nov 1990		5.195183	4.801335	5.589031	4.592252	5.798114
Dec 1990		5.276532	4.882106	5.670958	4.672717	5.880347
Jan 1991		5.086642	4.691602	5.481683	4.481888	5.691396
Feb 1991		5.020905	4.625148	5.416662	4.415055	5.626755

The above would be the forecast points in the original scale.

CATS

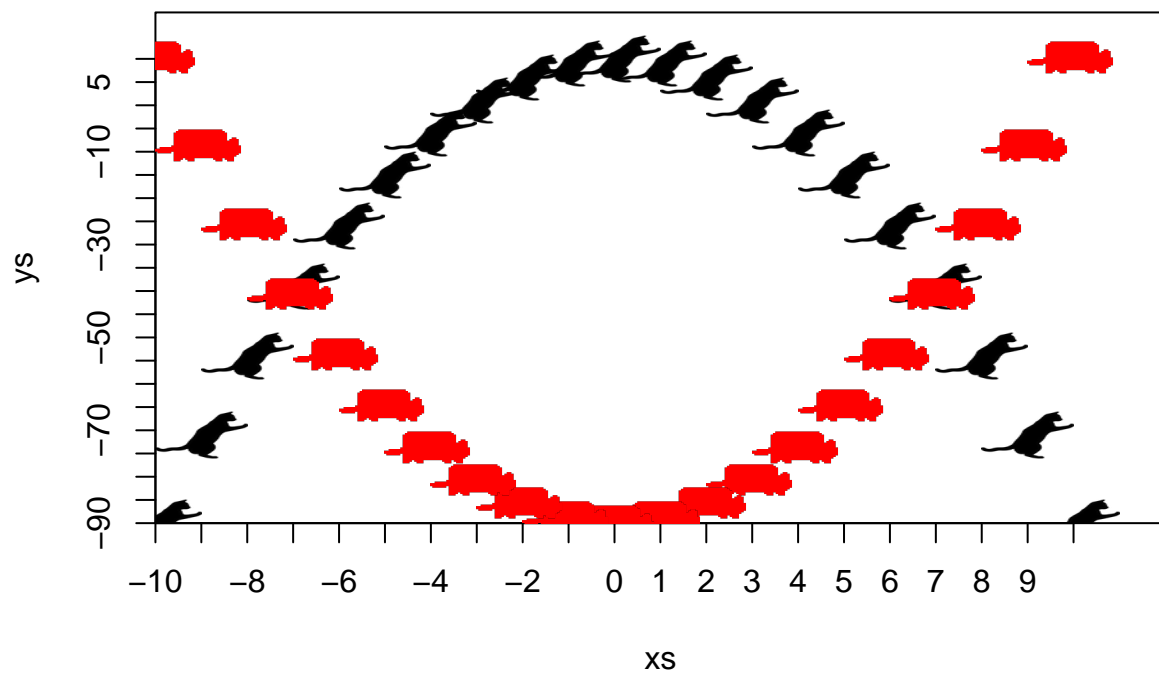
Final Note: Although this semester has been incredibly stressful and has given me anxiety to the point of not being able to sleep every night, and having such an accelerated heart rate that I was forced to start taking medication to prevent a cardiac arrest (along with daily exercise), and the material amount was ridiculous given the COVID19 situation... I am still very glad and happy that I took this course, and proud and satisfied with everything I learned, even after the tremendous amount of effort it costed me. That being said, I would like to thank Professor Steele for an amazing semester, and for caring about his students.

- PD1: It would be greatly appreciated if you could type your notex in Latex!
- PD2: Here are some random plots with cats instead of points $\hat{_}$

```
library(CatterPlots)
```

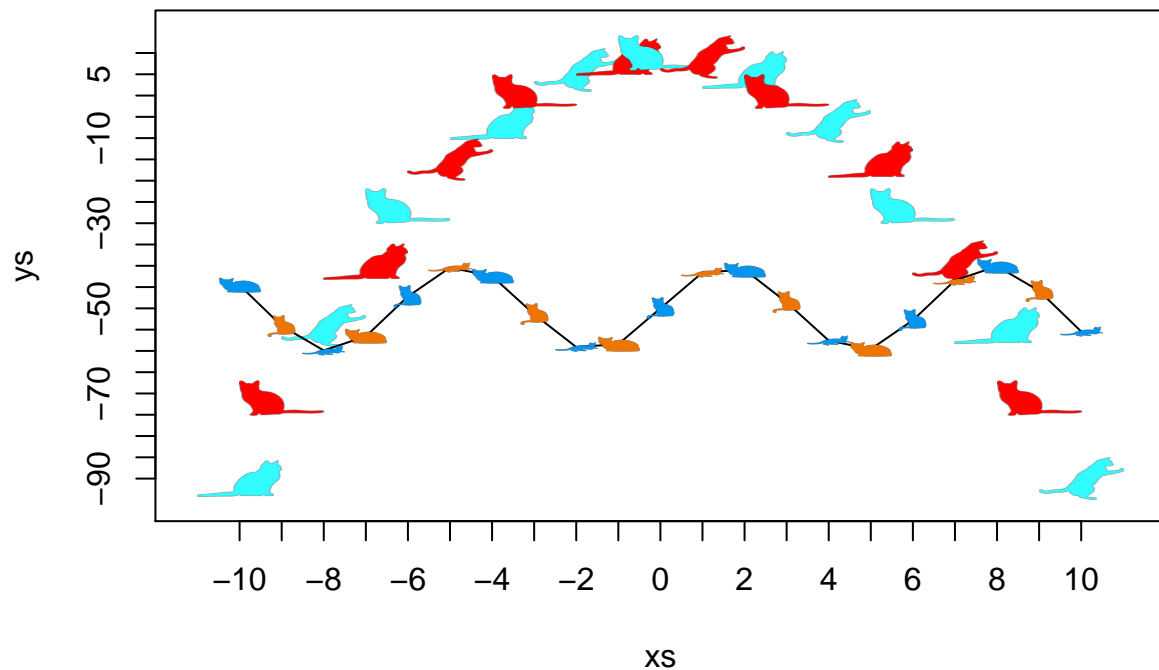
Welcome to CatterPlots.

```
x <- -10:10
y <- -x^2 + 10
purr <- catplot(xs=x, ys=y, cat=3, catcolor='#000000FF')
cats(purr, -x, -y, cat=11, catcolor='#FF0000')
```



```
# for more fun ...
```

```
meow <- multicat(xs=x, ys=y, cat=c(1,2,3), catcolor=list('#33FCFF','#FF0000'), canvas=c(-0.1,1.1, -0.1,
morecats(meow, x, 10*sin(x)+40, size=0.05, cat=c(4,5,6), catcolor=list('#0495EE','#EE7504'), type="line
```



```
-0.025  0.429402111088937  0.025  0.479402111088937
0.025  0.333788151475824  0.075  0.383788151475824
0.075  0.276064175337662  0.125  0.326064175337662
0.125  0.309301340128121  0.175  0.359301340128121
0.175  0.402941549819893  0.225  0.452941549819893
0.225  0.470892427466314  0.275  0.520892427466314
0.275  0.450680249530793  0.325  0.500680249530793
0.325  0.360887999194013  0.375  0.410887999194013
0.375  0.284070257317432  0.425  0.334070257317432
0.425  0.29085290151921  0.475  0.34085290151921
0.475  0.375  0.525  0.425
0.525  0.45914709848079  0.575  0.50914709848079
0.575  0.465929742682568  0.625  0.515929742682568
0.625  0.389112000805987  0.675  0.439112000805987
0.675  0.299319750469207  0.725  0.349319750469207
0.725  0.279107572533686  0.775  0.329107572533686
0.775  0.347058450180107  0.825  0.397058450180107
0.825  0.440698659871879  0.875  0.490698659871879
0.875  0.473935824662338  0.925  0.523935824662338
0.925  0.416211848524176  0.975  0.466211848524176
0.975  0.320597888911063  1.025  0.370597888911063
```

```
# random cats
meow <- multicat(xs=x, ys=rnorm(21),
                 cat=c(1,2,3,4,5,6,7,8,9,10),
```

```
catcolor=list('#33FCFF'),  
canvas=c(-0.1,1.1, -0.1, 1.1),  
xlab="some cats", ylab="other cats", main="Random Cats")
```

